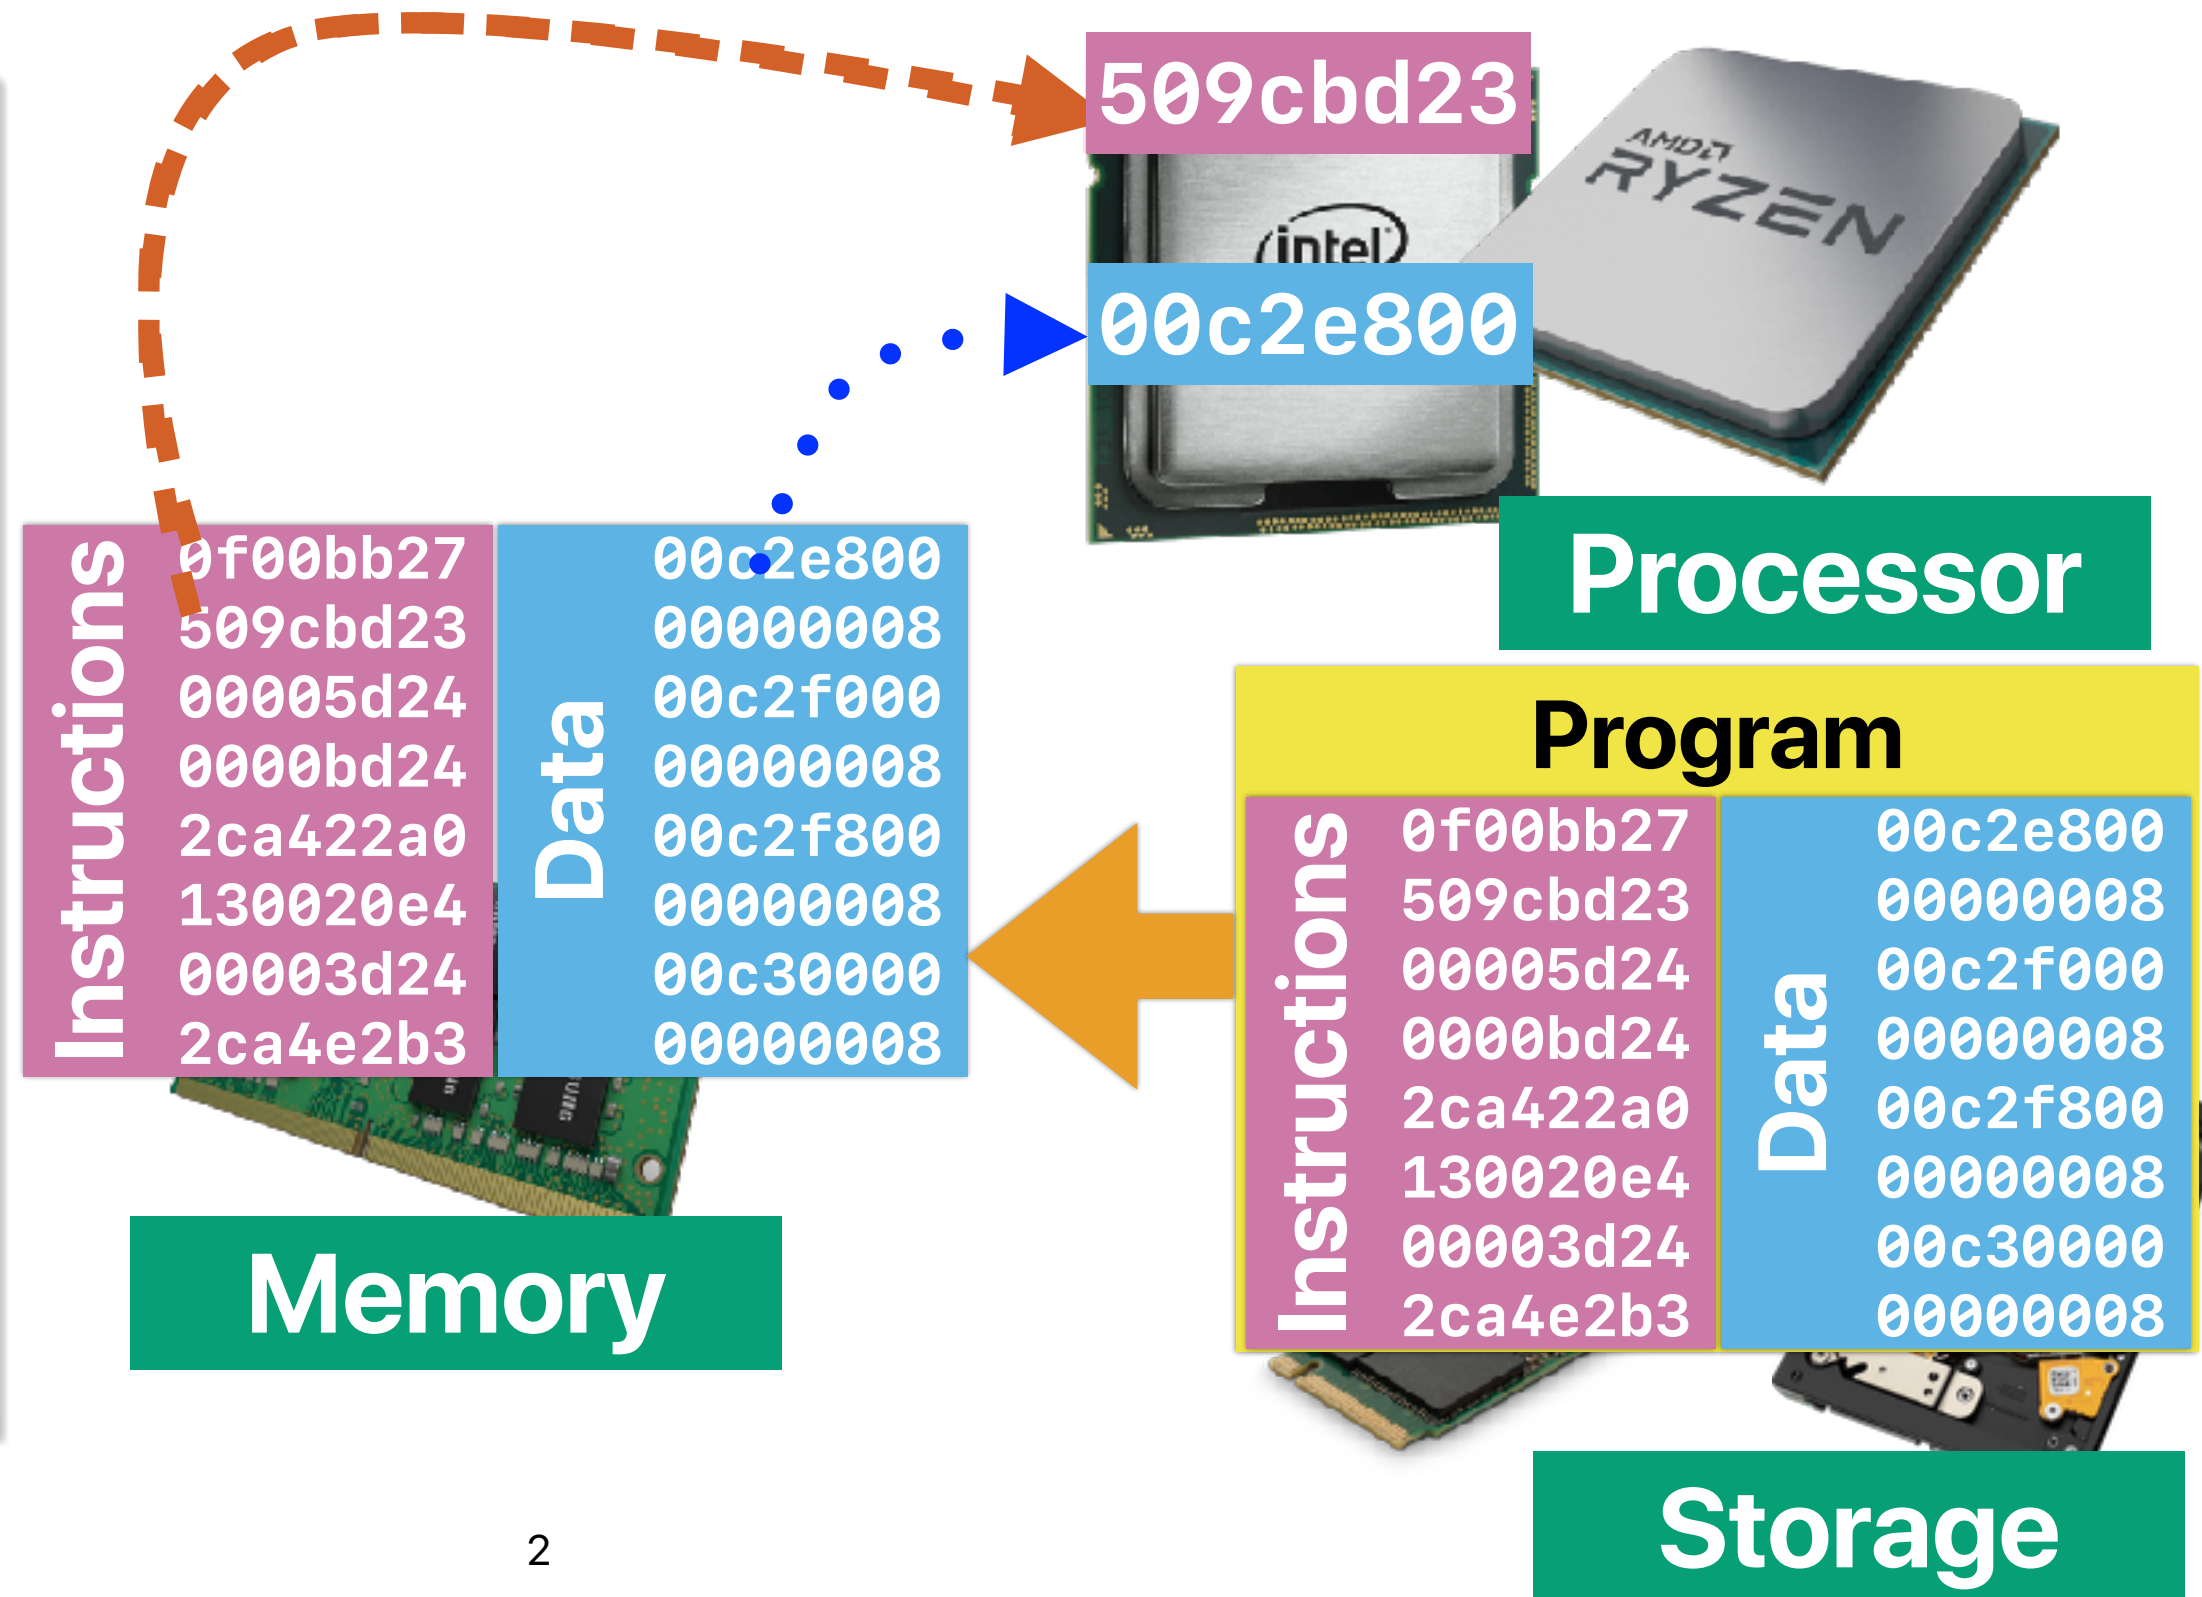


Modern Processor Design (I): in the pipeline

Hung-Wei Tseng

von Neumann Architecture



Recap: Microprocessor — a collection of functional units



Instructions

Instruction Set Architecture

Logical
operations

Simple
Arithmetic
Operations
(Add/Sub)

Complex
Arithmetic
Operations
(Mul/div)

Branch/
Jump

Memory
Operations

Processor

Tricky C/C++ programming questions?

- Give a fastest way to multiply any number by 9
- How to measure the size of any variable without "sizeof" operator?.
- How to measure the size of any variable without using "sizeof" operator?
- Write code snippets to swap two variables in five different ways
- How to swap between first & 2nd byte of an integer in one line statement?
- What is the efficient way to divide a no. by 4?
- Suggest an efficient method to count the no. of 1's in a 32 bit no. Remember without using loop & testing each bit.
- Test whether a no. is power of 2 or not.
- How to check endianness of the computer.
- Write a C-program which does the addition of two integers without using '+' operator.
- Write a C-program to find the smallest of three integers without using any of the comparison operators.
- Find the maximum & minimum of two numbers in a single line without using any condition & loop.
- What "condition" expression can be used so that the following code snippet will print Hello world.
- How to print number from 1 to 100 without using conditional operators.
- WAP to print 100 times "Hello" without using loop & goto statement.
- Write the equivalent expression for $x\%8$.

<https://www.emblogic.com/blog/12/tricky-c-interview-questions/>

Recap: Demo (3) — Bitwise operations?

using bit-shifting

d. /* one line statement using bit-wise operators */ (most efficient)

```
a^=b^=a^=b;
```

The order of evaluation is from right to left. This is same as in approach (c) but the three statements are compounded into one statement.

A

```
void regswap(int* a, int* b) {  
    int temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

B

```
void xorswap(int* a, int* b) {  
    *a ^= *b ^= *a = *b;  
}
```

Recap: Leveraging more “bit-wise” operations in C code will make the program significantly faster



Recap: Why adding a sort makes it faster

- Why the sorting the array speed up the code despite the increased instruction count?

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize; ++i) {
        if (data[i] >= threshold)
            sum ++;
    }
}
```

Outline

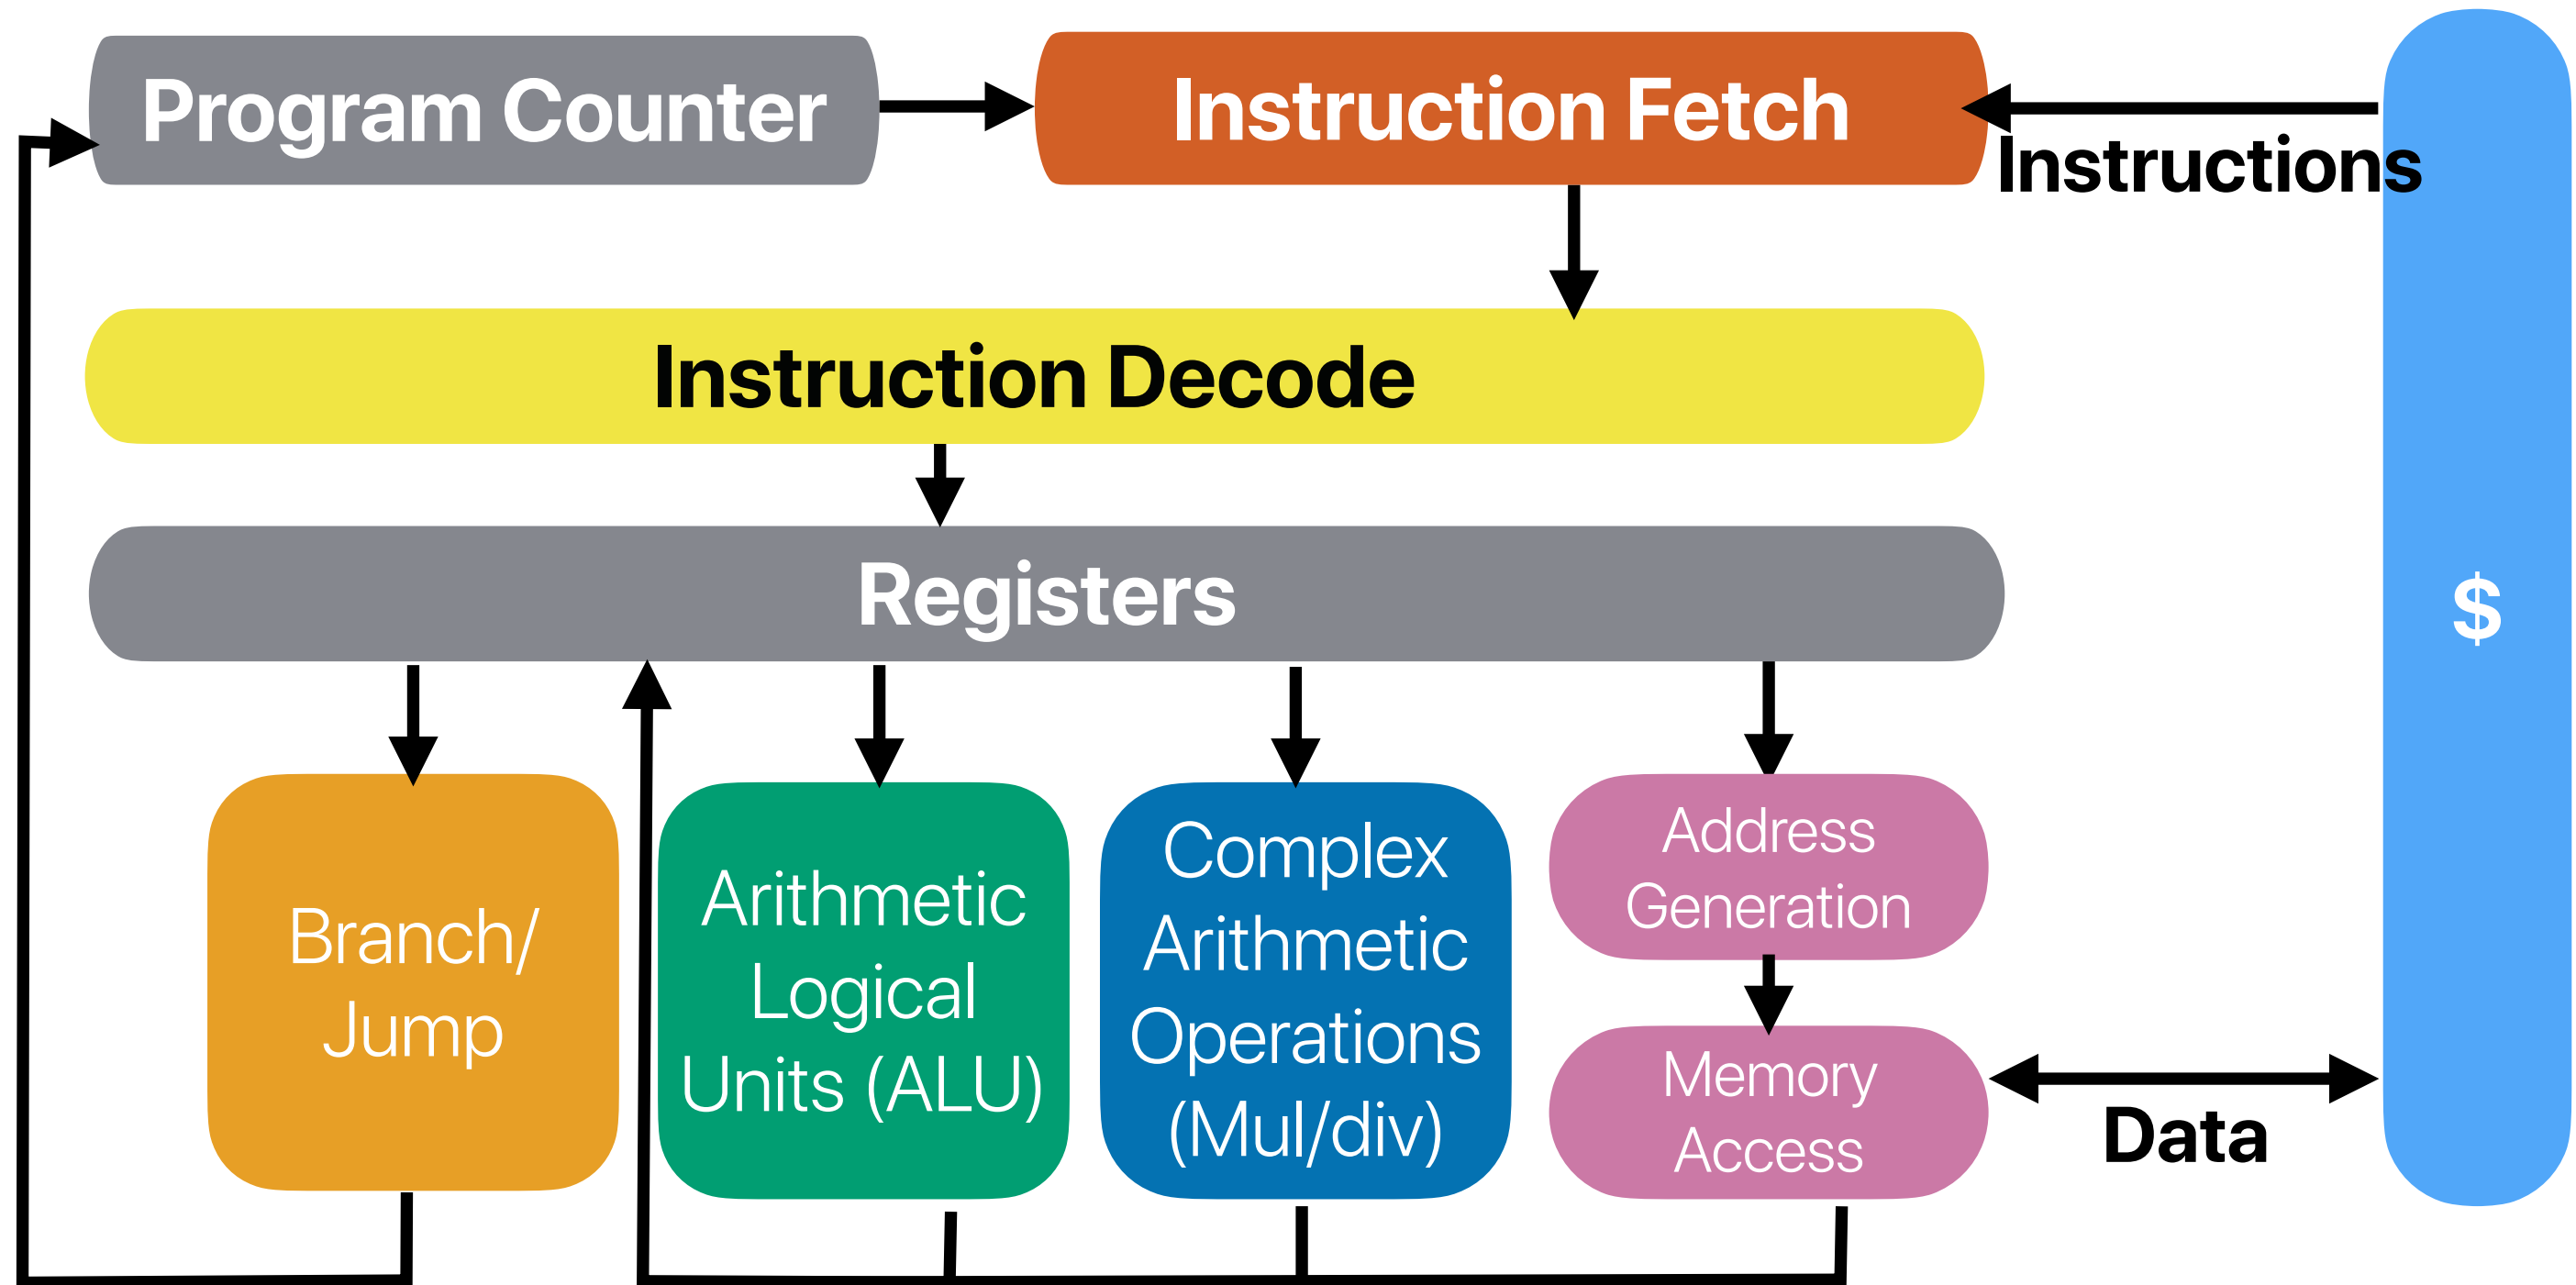
- Recap: the concept of a processor
- Pipelined Processor
- Pipeline Hazards
 - Structural Hazards
 - Control Hazards
 - Data Hazards

Basic Processor Design

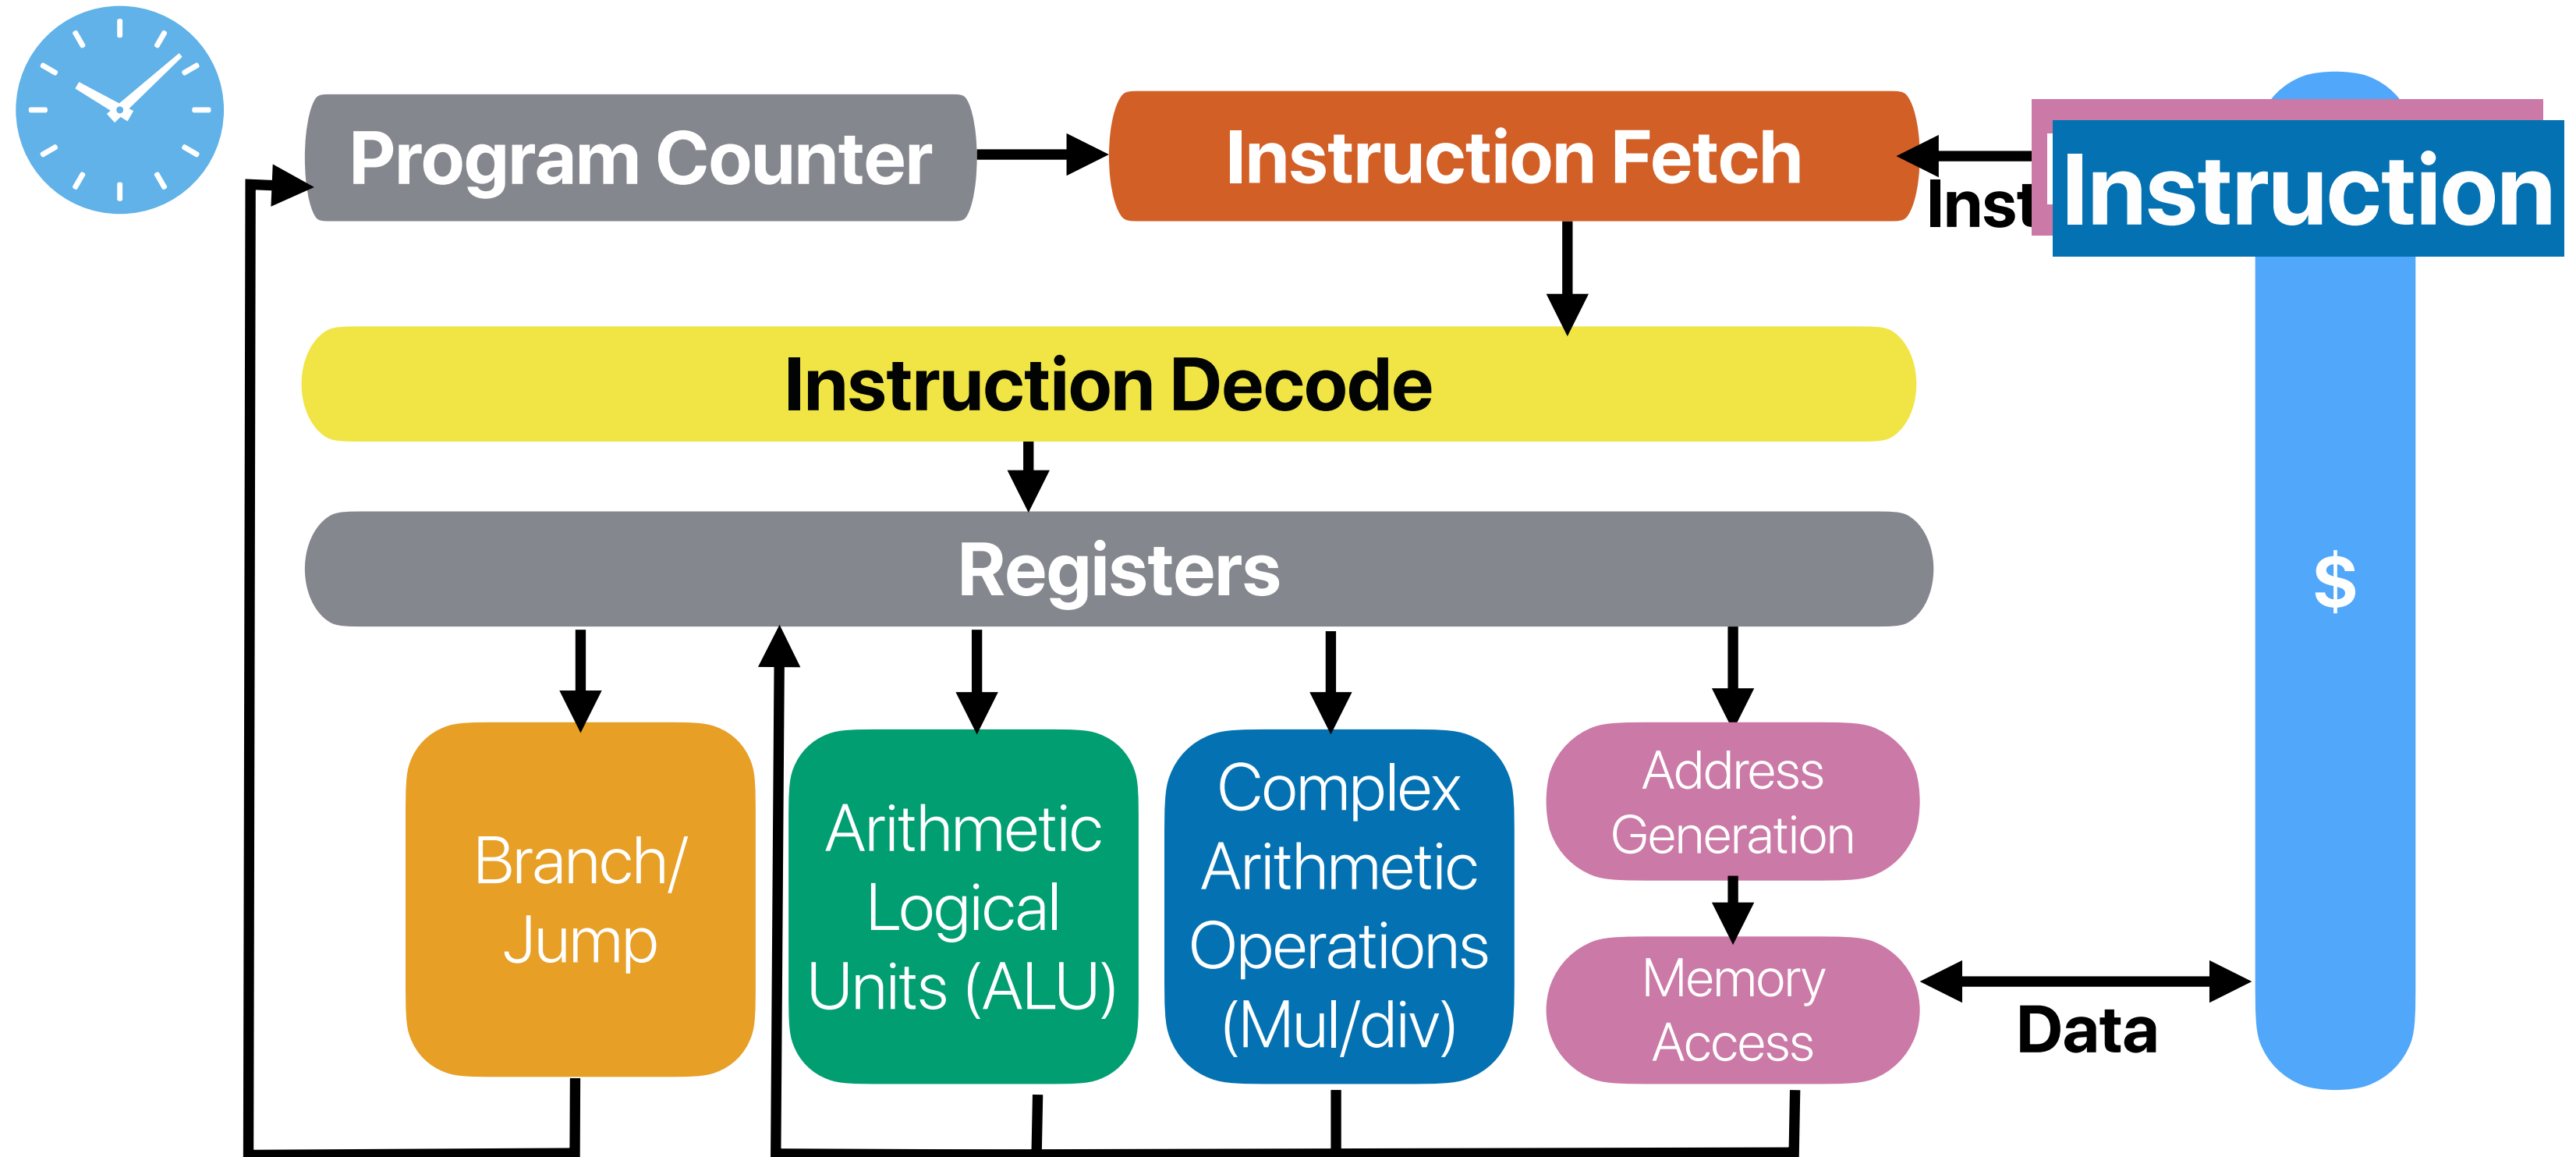
The “life” of an instruction

- Instruction Fetch (**IF**) — fetch the instruction from memory
- Instruction Decode (**ID**)
 - Decode the instruction for the desired operation and operands
 - Reading source register values
- Execution (**EX**)
 - ALU instructions: Perform ALU operations
 - Conditional Branch: Determine the branch outcome (taken/not taken)
 - Memory instructions: Determine the effective address for data memory access
- Data Memory Access (**MEM**) — Read/write memory
- Write Back (**WB**) — Present ALU result/read value in the target register
- Update PC
 - If the branch is taken — set to the branch target address
 - Otherwise — advance to the next instruction — current PC + 4


Functional Units of a Microprocessor



If we want to perform one instruction each cycle...



Simple implementation w/o branch




Pipelining



Pipelining

- Different parts of the processor works on different instructions simultaneously
- A processor is now working on multiple instructions from the same program (though on different stages) simultaneously.
 - **ILP: Instruction-level parallelism**
- A **clock** signal controls and synchronize the beginning and the end of each part of the work
- A **pipeline register** between different parts of the processor to keep intermediate results necessary for the upcoming work

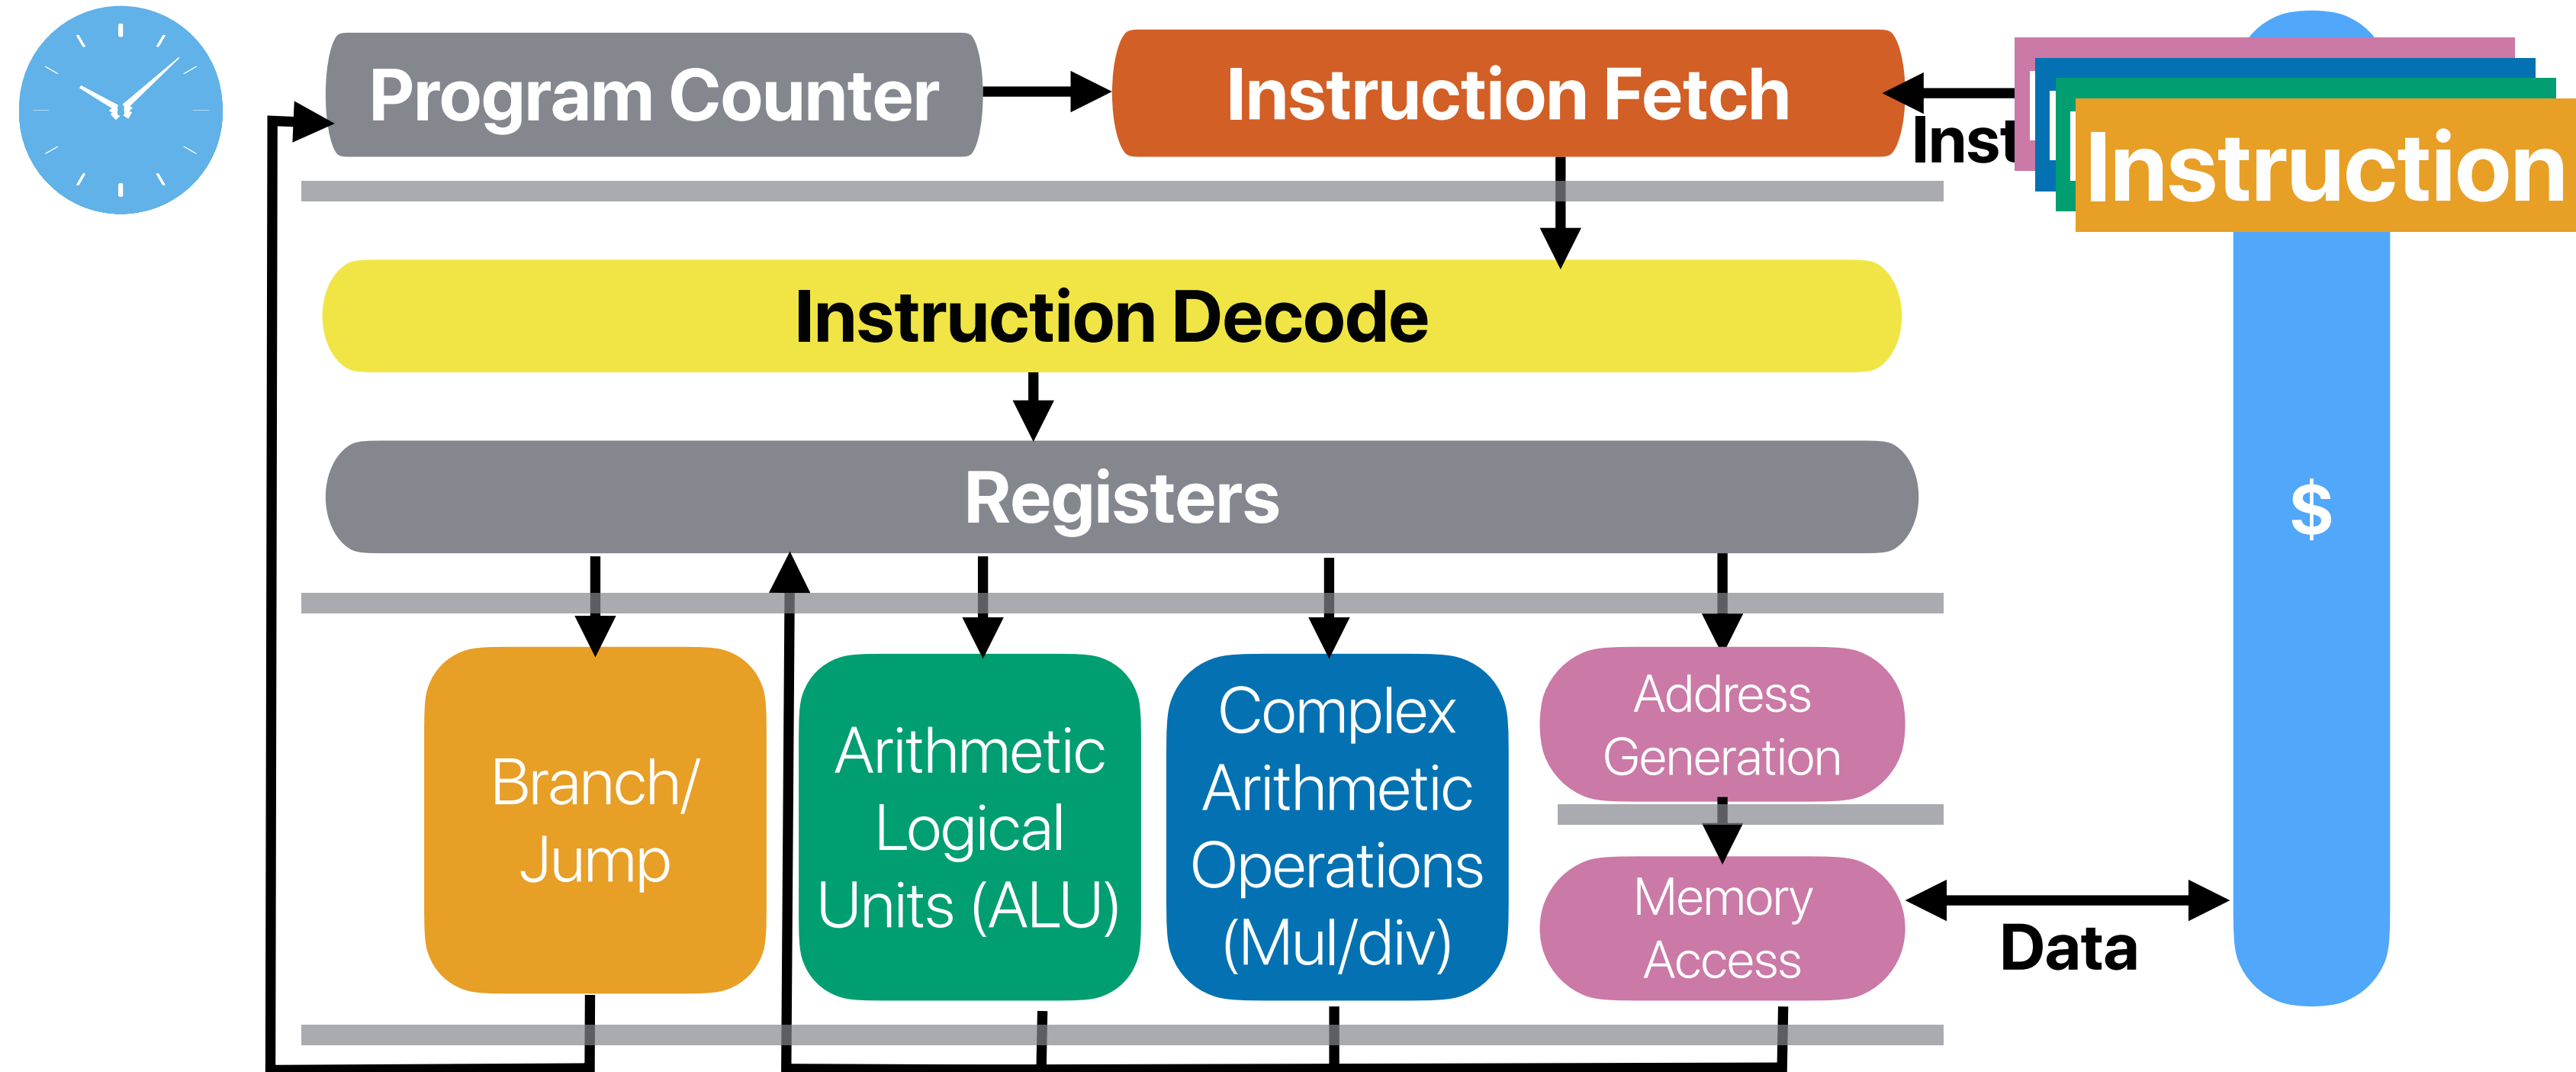
Pipelining



Pipelining

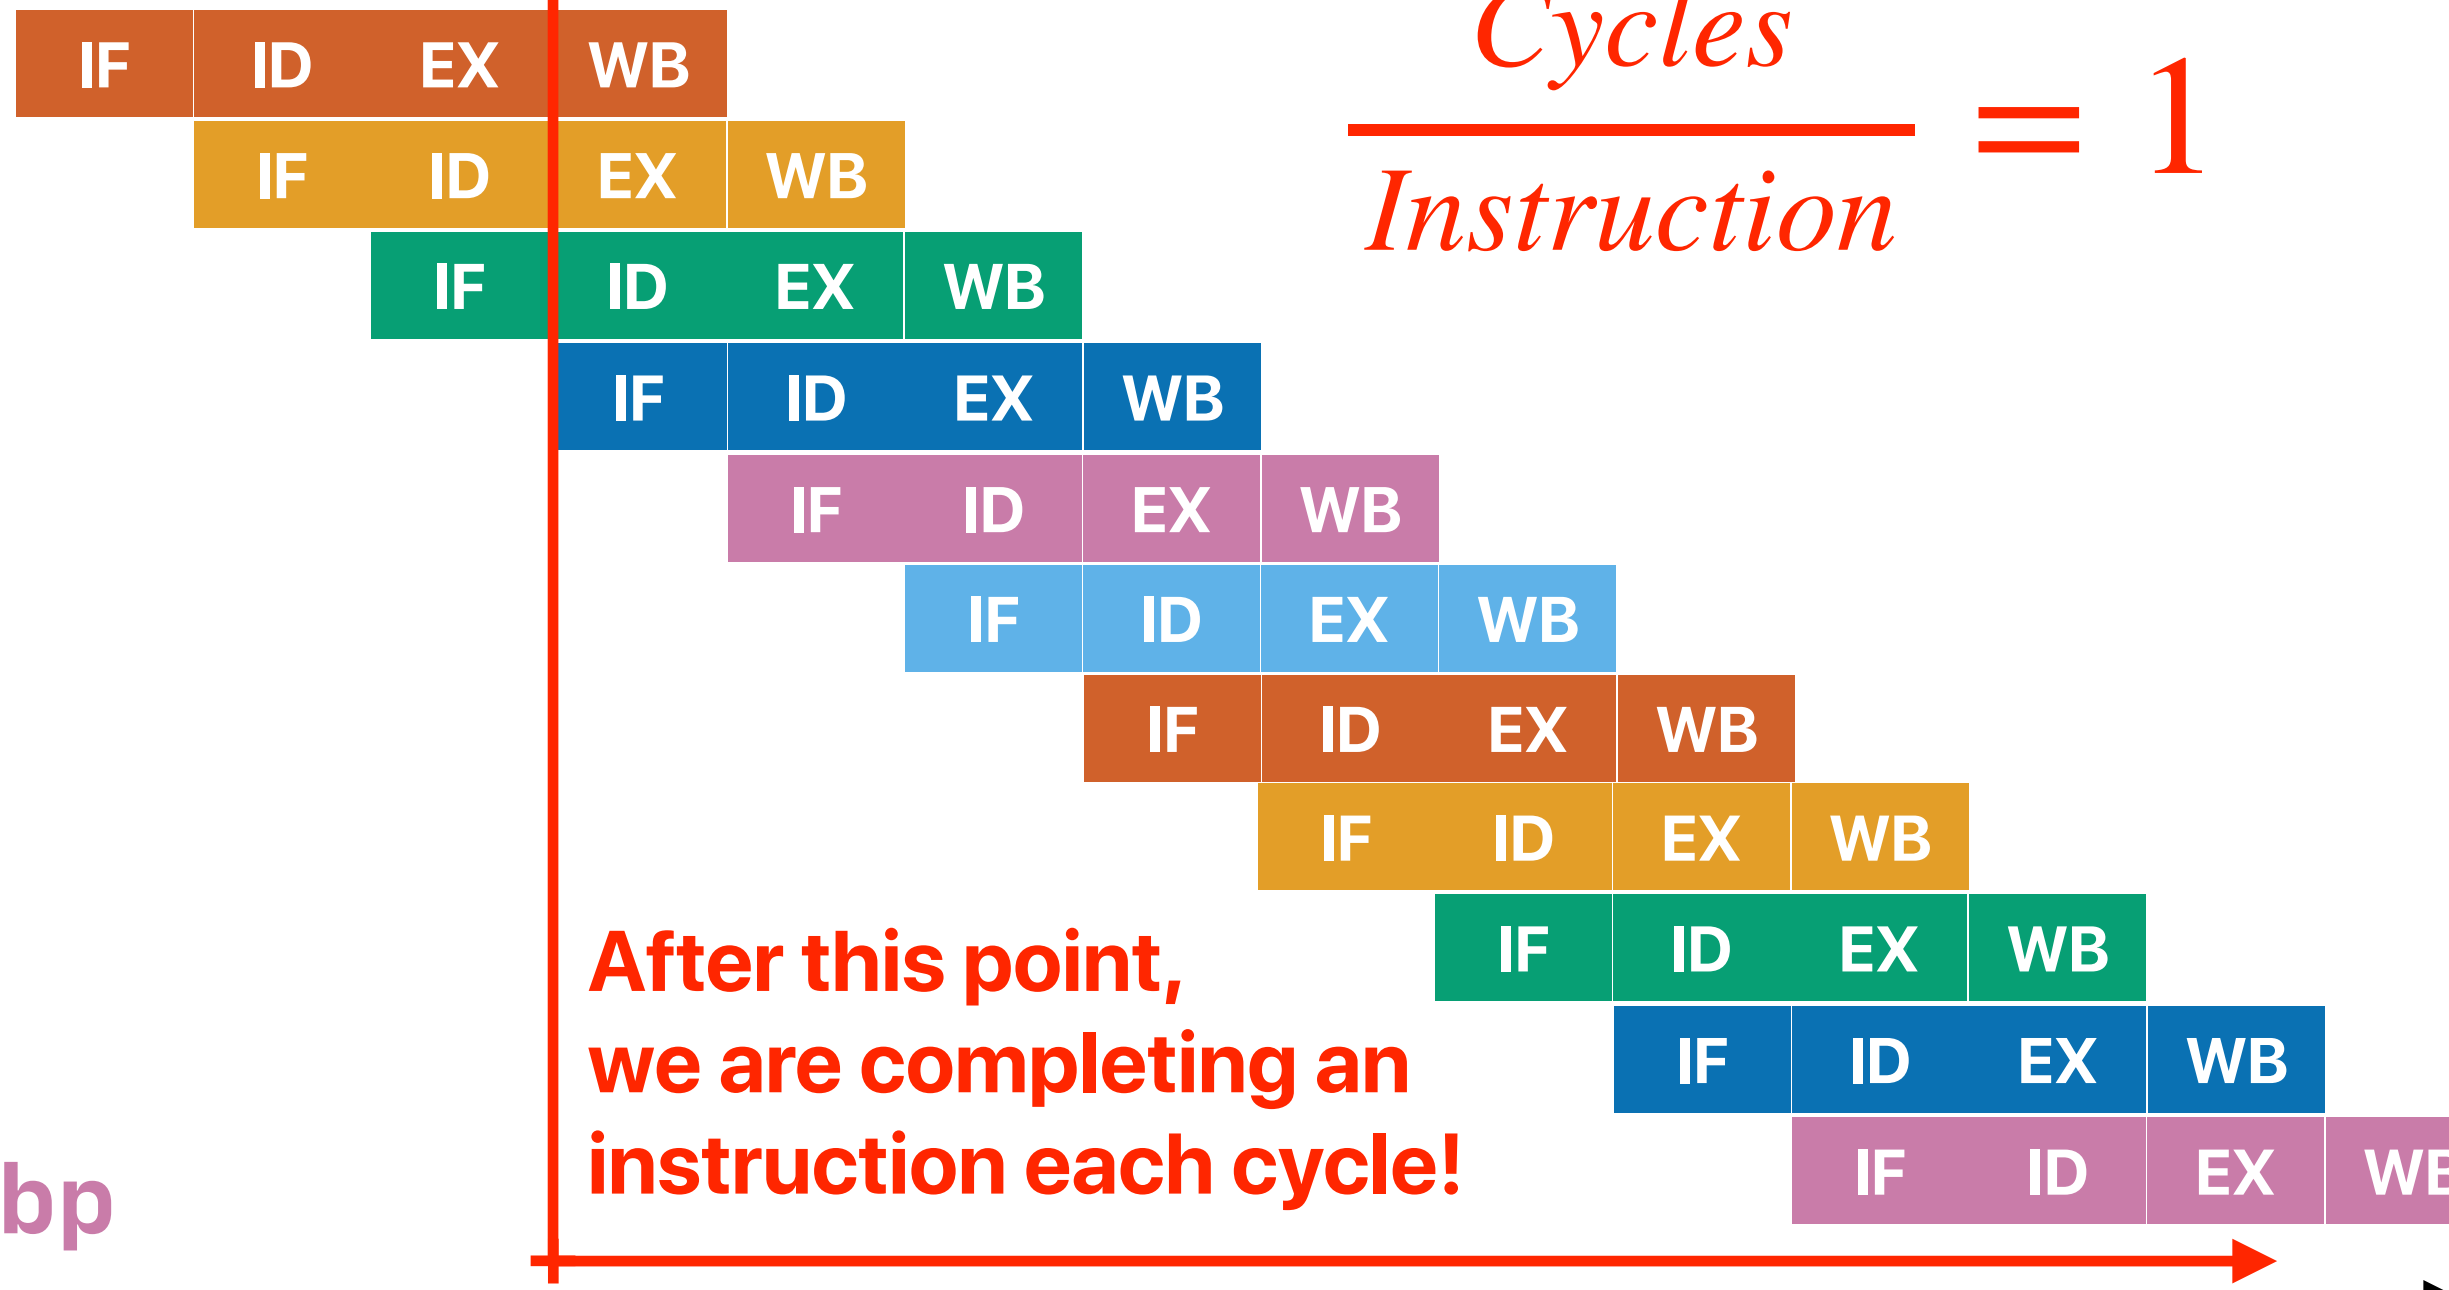


Pipelined execution



Pipelining

```
addl    %eax, %eax
addl    %rdi, %ecx
addq    $4, %r11
testl   %esi, %esi
movl    $10, %edx
pushq   %r12
pushq   %rbp
pushq   %rbx
subq    $8, %rsp
addl    %rsi, %rdi
movslq  %eax, %rbp
```

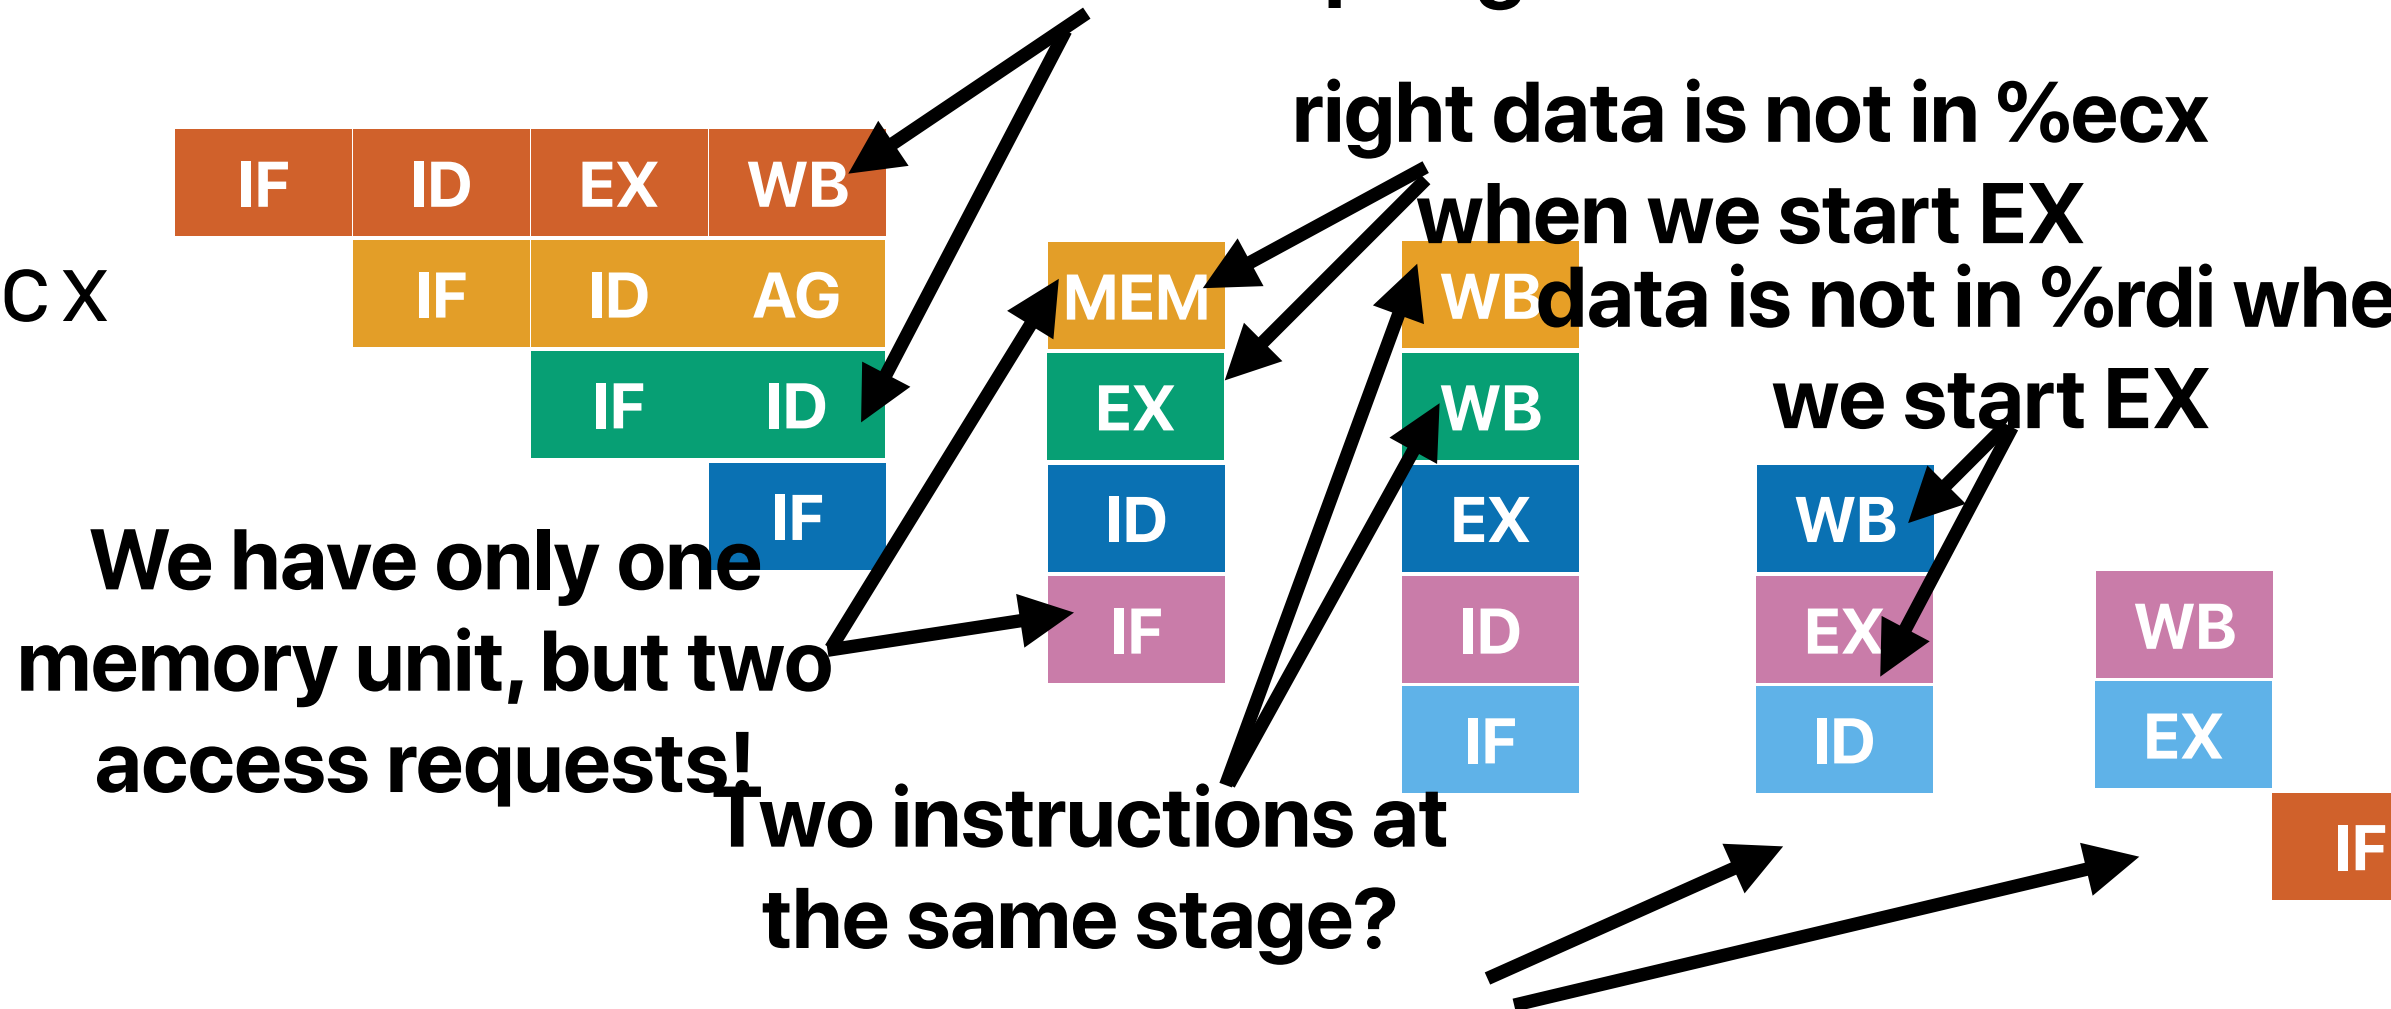


$$\frac{\text{Cycles}}{\text{Instruction}} = 1$$

After this point,
we are completing an
instruction each cycle!

Pipelining

① xorl %eax, %eax
② movl (%rdi), %ecx
③ addl %ecx, %eax
④ addq \$4, %rdi
⑤ cmpq %rdx, %rdi
⑥ jne .L3
⑦ ret



Takeaways: pipeline processors

- Pipelining helps to improve the throughput of processors
 - Allowing shorter cycle time as each cycle only make progress for part of an instruction
 - Different pipeline stages work on different instructions concurrently
 - Theoretical CPI remains the same as single-cycle design and the throughput/speedup is in proportion to the speedup of cycle time

Pipeline hazards

Three types of pipeline hazards

- Structural hazards — resource conflicts cannot support simultaneous execution of instructions in the pipeline
- Control hazards — the PC can be changed by an instruction in the pipeline
- Data hazards — an instruction depending on a the result that's not yet generated or propagated when the instruction needs that


Takeaways: pipeline processors

- Pipelining helps to improve the throughput of processors
 - Allowing shorter cycle time as each cycle only make progress for part of an instruction
 - Different pipeline stages work on different instructions concurrently
 - Theoretical CPI remains the same as single-cycle design and the throughput/speedup is in proportion to the speedup of cycle time
- Pipeline hazards prevent us from reaching the theoretical CPI
 - Structural hazards
 - Control hazards
 - Data hazards

**Stall — the universal solution to
pipeline hazards**

Stall whenever we have a hazard

- Stall: the hardware allows the earlier instruction to proceed, all later instructions stay at the same stage
- Disable the pipeline register update for later instructions
- The stalled instructions still have the same input from the pipeline registers



Slow! — 4 additional cycles

Structural Hazards

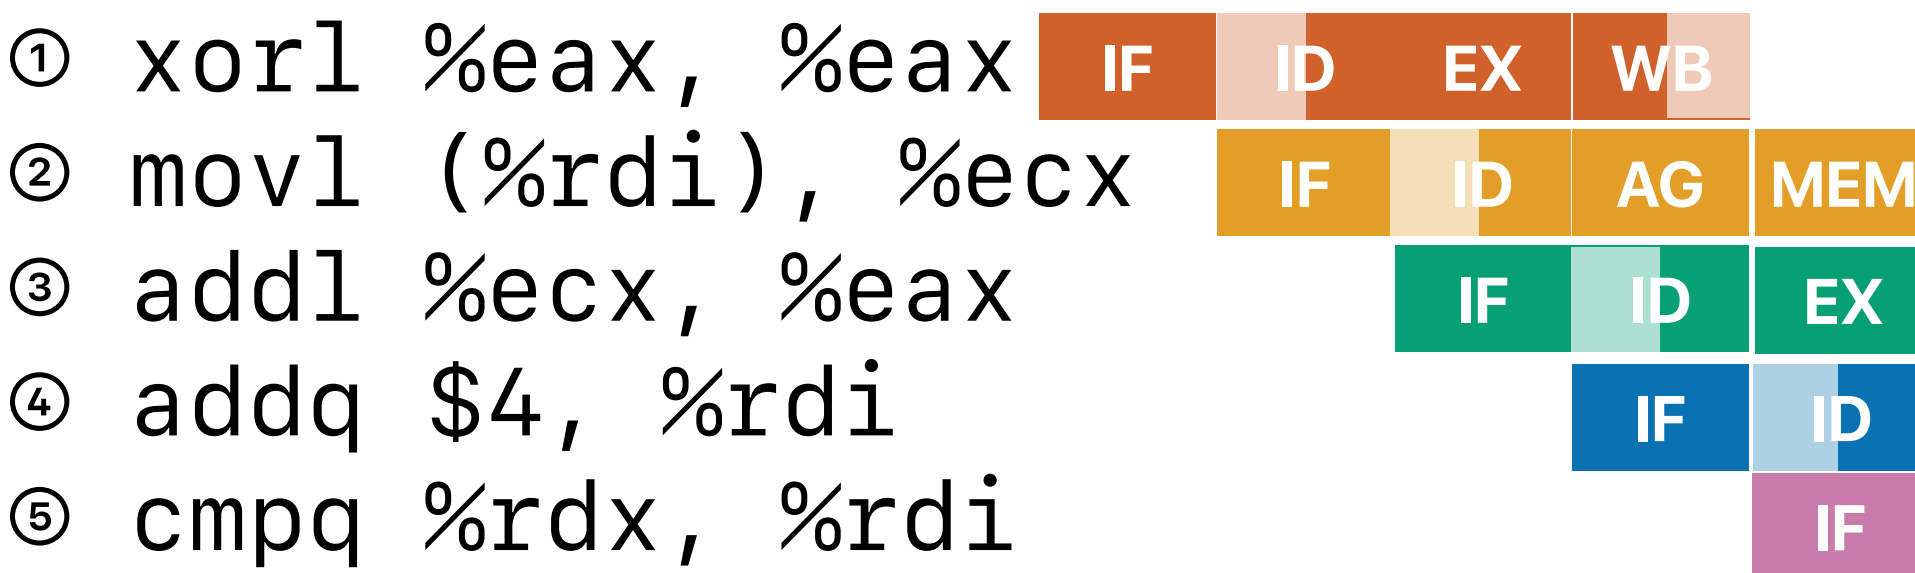
Dealing with the conflicts between ID/WB

- The same register cannot be read/written at the same cycle
- Better solution: write early, read late
 - Writes occur at the clock edge and complete long enough before the end of the clock cycle.
 - This leaves enough time for outputs to settle for reads
 - The revised register file is the default one from now!

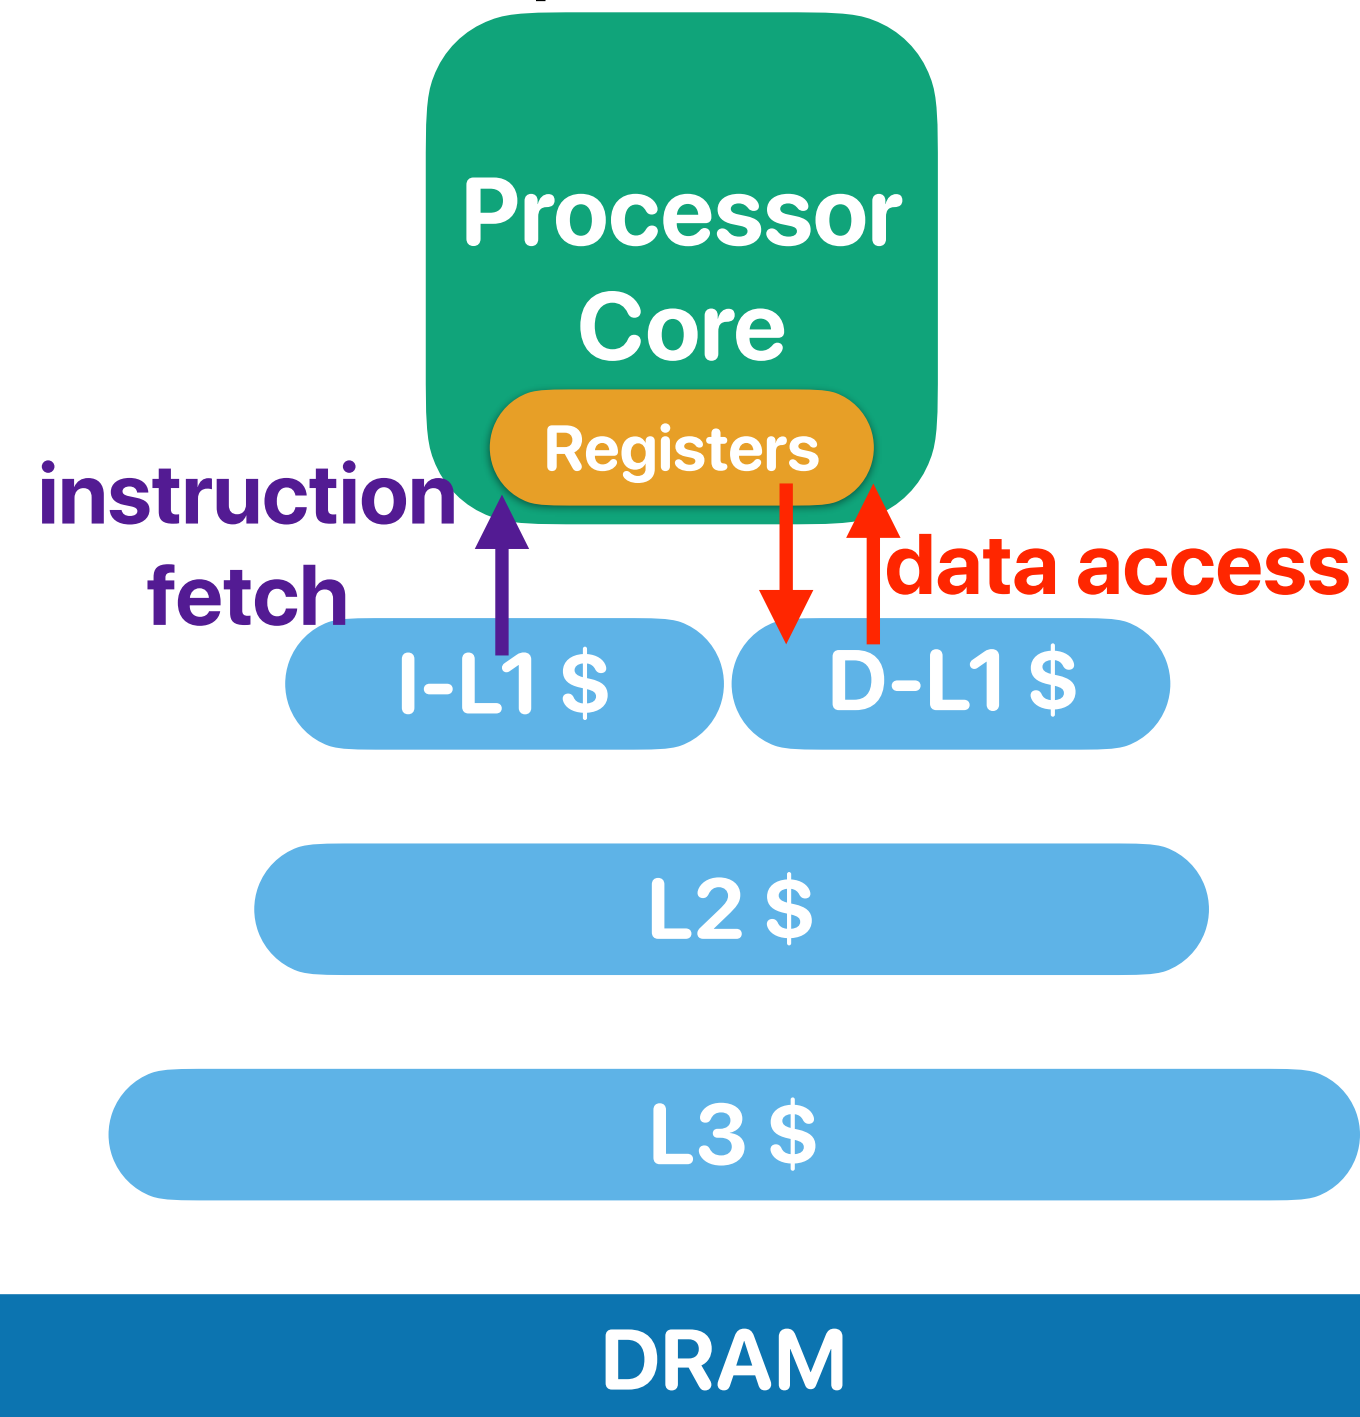


How to with the conflicts between MEM and IF?


- The memory unit can only accept/perform one request each cycle



"Split L1" cache!




Split L1-\$



Both (2) and (3) want to "WB"


- The memory unit can only accept/perform one request each cycle



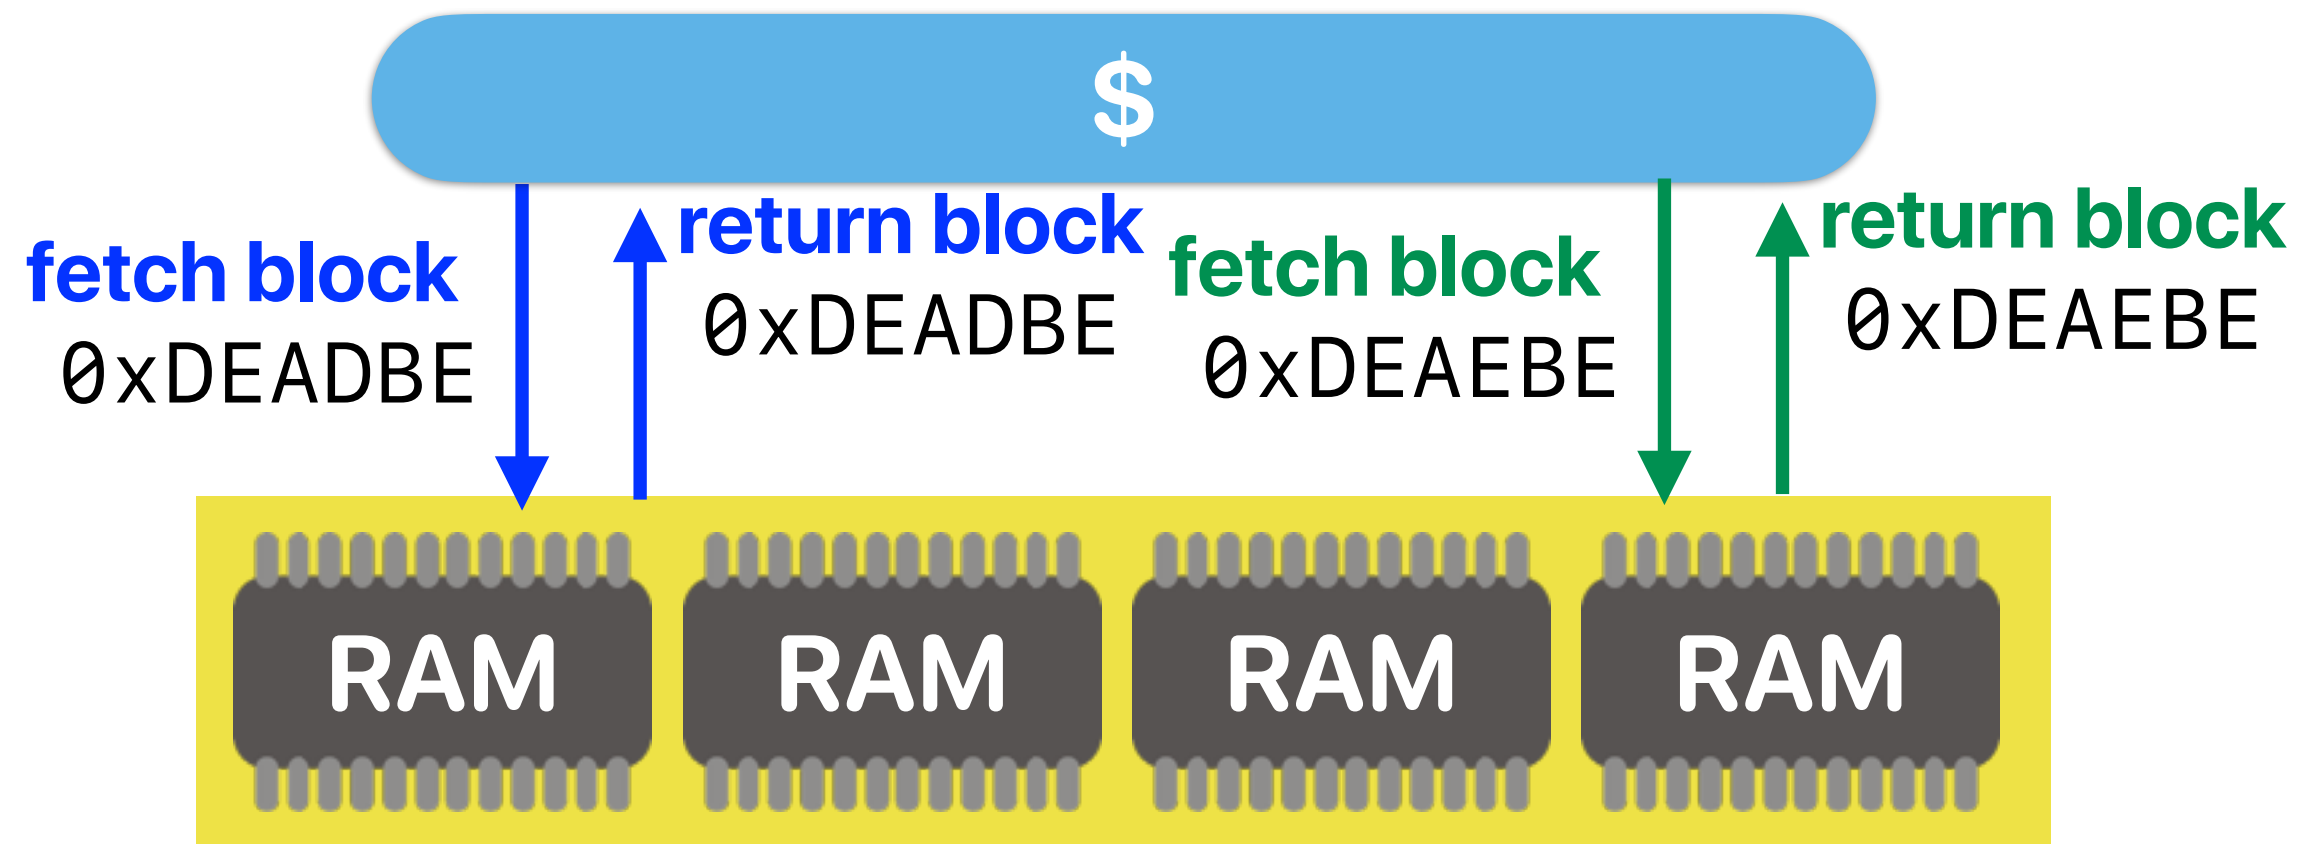
(3) has to stall

What if we've back-to-back memory instructions?

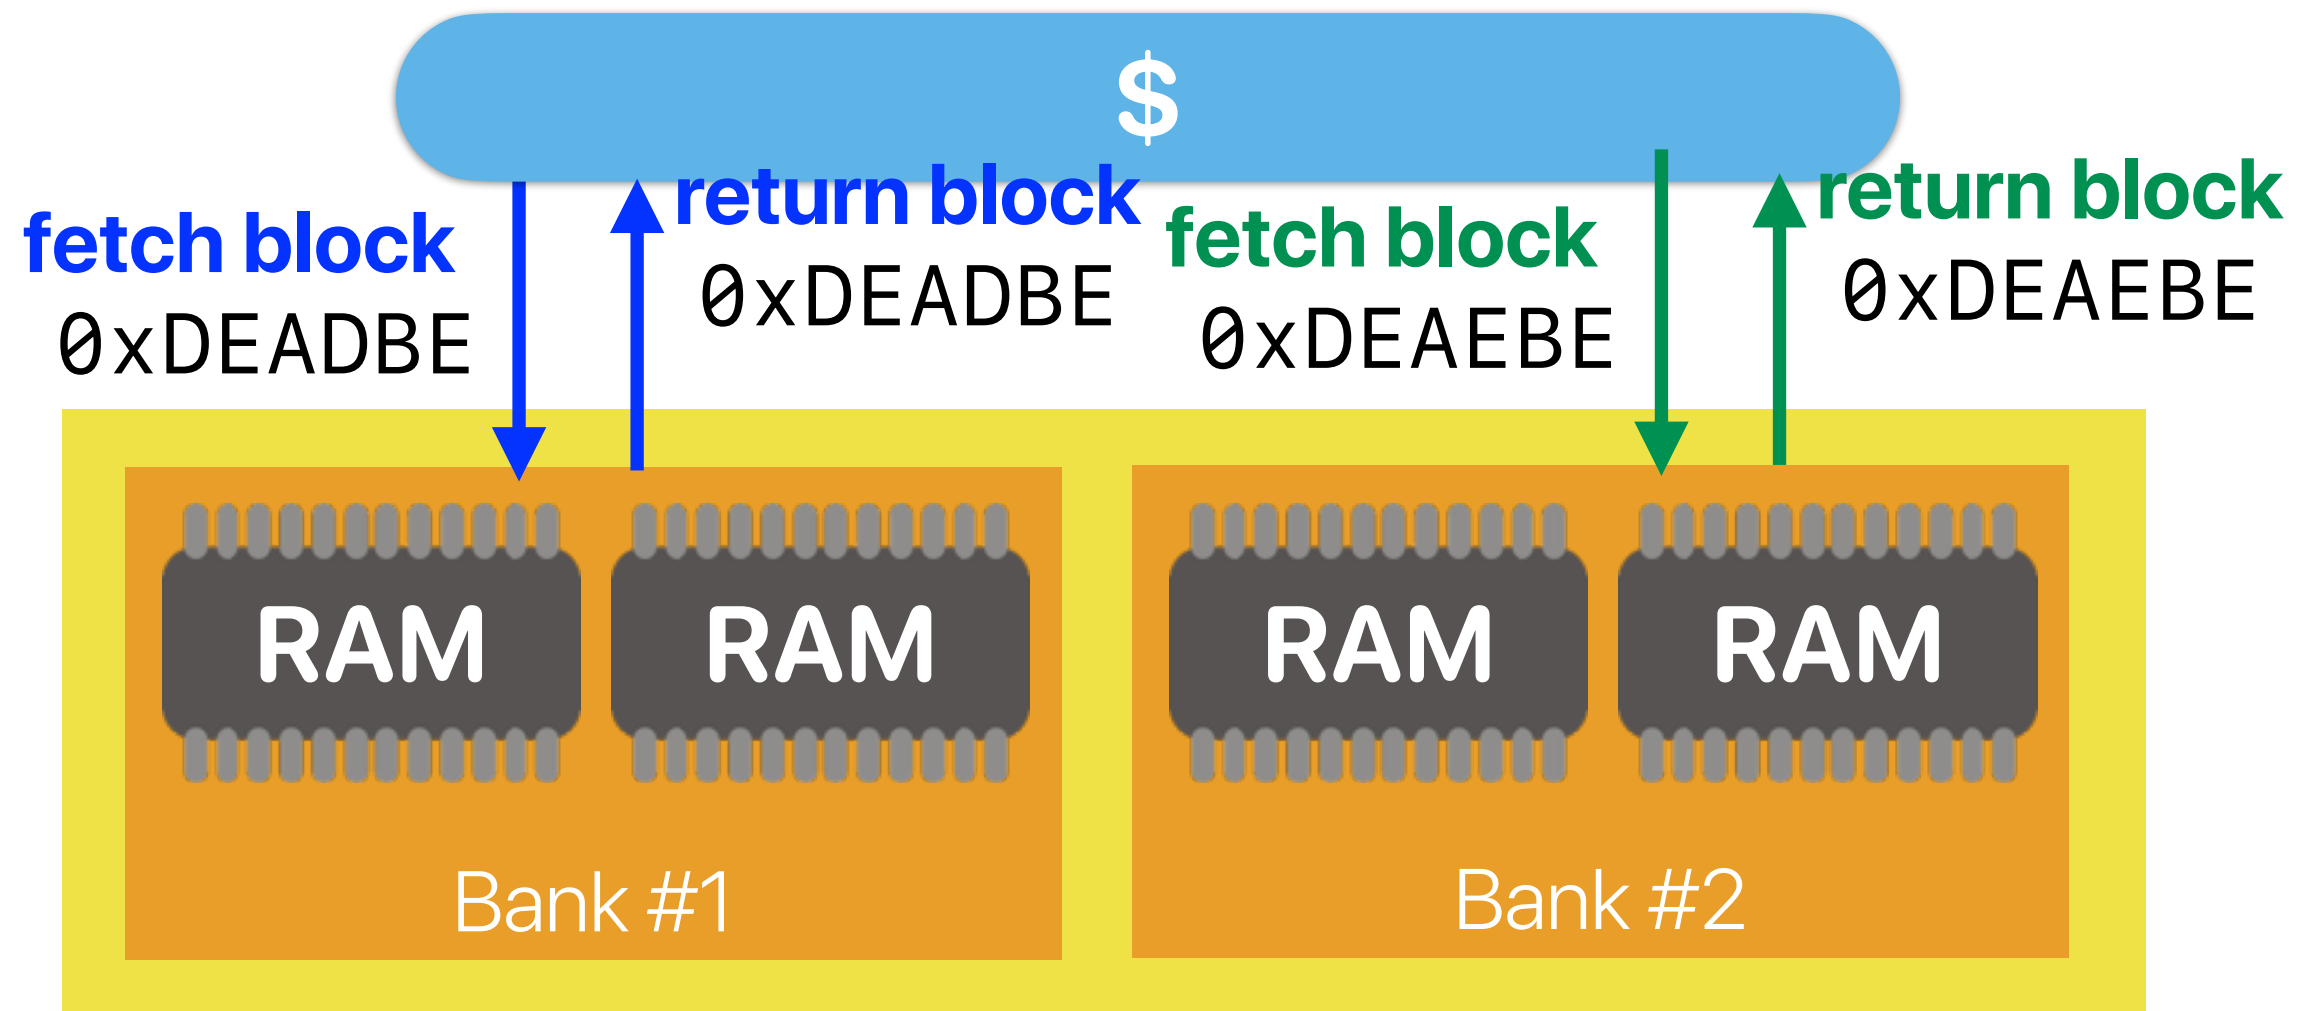
- Every instruction needs to go through exactly the same number of stages



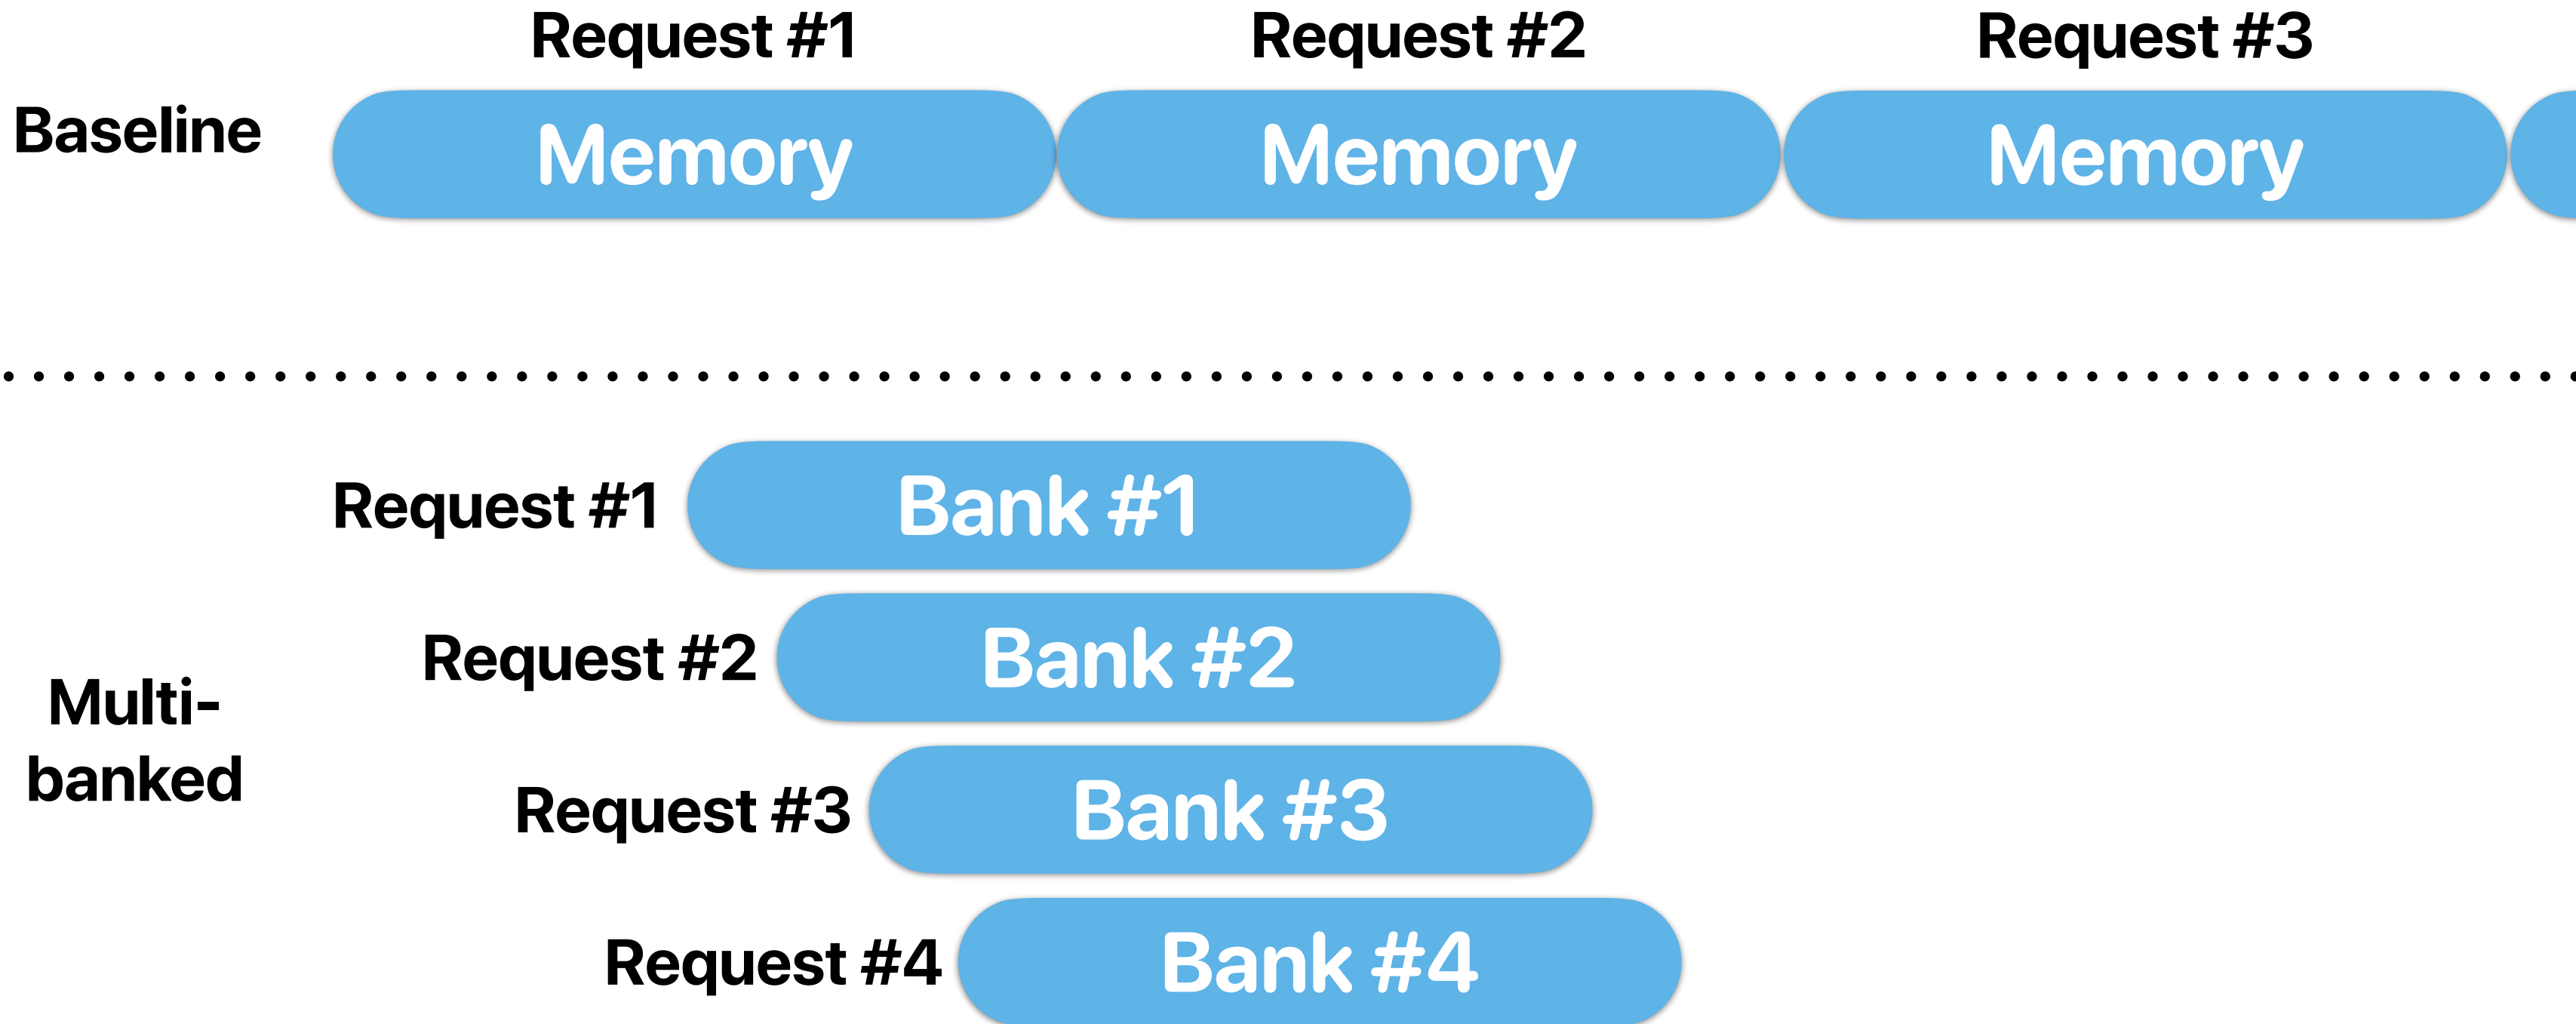
Blocking cache



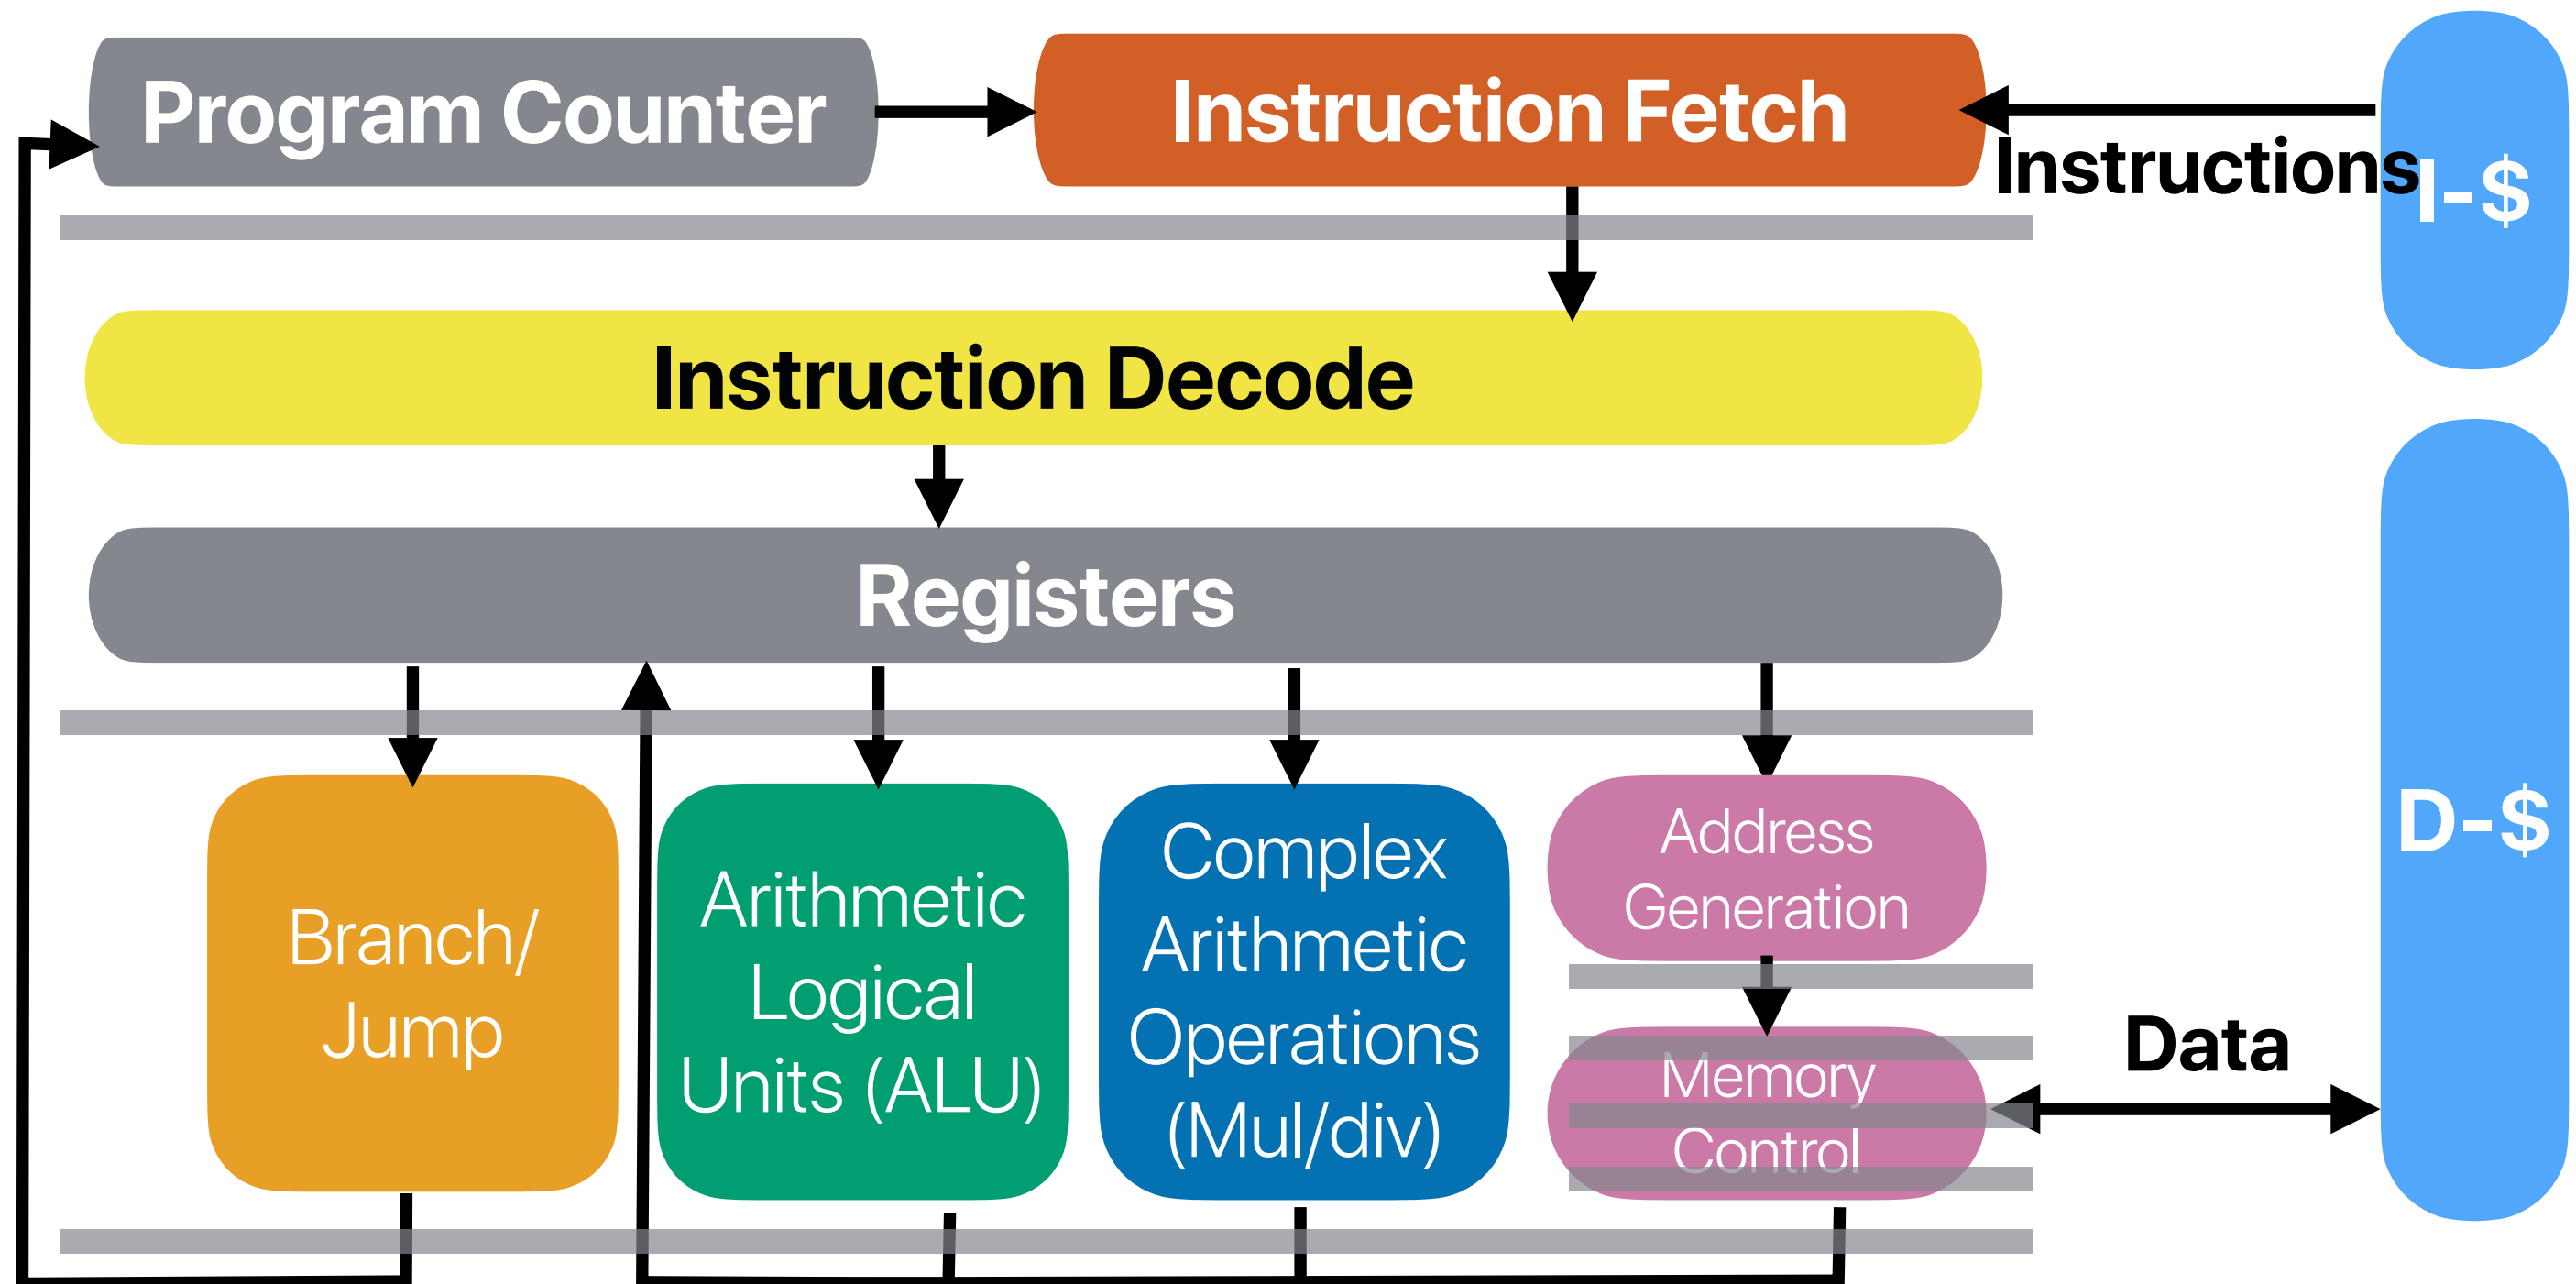
Multibanks & non-blocking caches



Pipelined access and multi-banked caches



Split L1-\$ + Multibanked, non-blocking cache



Takeaways: pipeline processors

- Pipelining helps to improve the throughput of processors
 - Allowing shorter cycle time as each cycle only make progress for part of an instruction
 - Different pipeline stages work on different instructions concurrently
 - Theoretical CPI remains the same as single-cycle design and the throughput/speedup is in proportion to the speedup of cycle time
- Pipeline hazards prevent us from reaching the theoretical CPI
 - Structural hazards
 - Control hazards
 - Data hazards
- The most efficient approach to address structural hazards is to make the hardware available to support concurrent execution
 - Register file
 - Split caches
 - Multibanked, non-blocking cache