

# **Modern Processor Design (III): Whenever You're Ready**

Hung-Wei Tseng

# Recap: how does compiler implement if-else?

```
for(i=0;i<size;i++)
{
    if (data[i] % 2 != 0)
        call_when_true(&data[i]);
    else
        call_when_false(&data[i]);
}
return 0;
```

go to L10 when data[i] % 2 == 1

if (data[i] % 2 != 0) // branch taken if it's false

call\_when\_true(&data[i]);

else

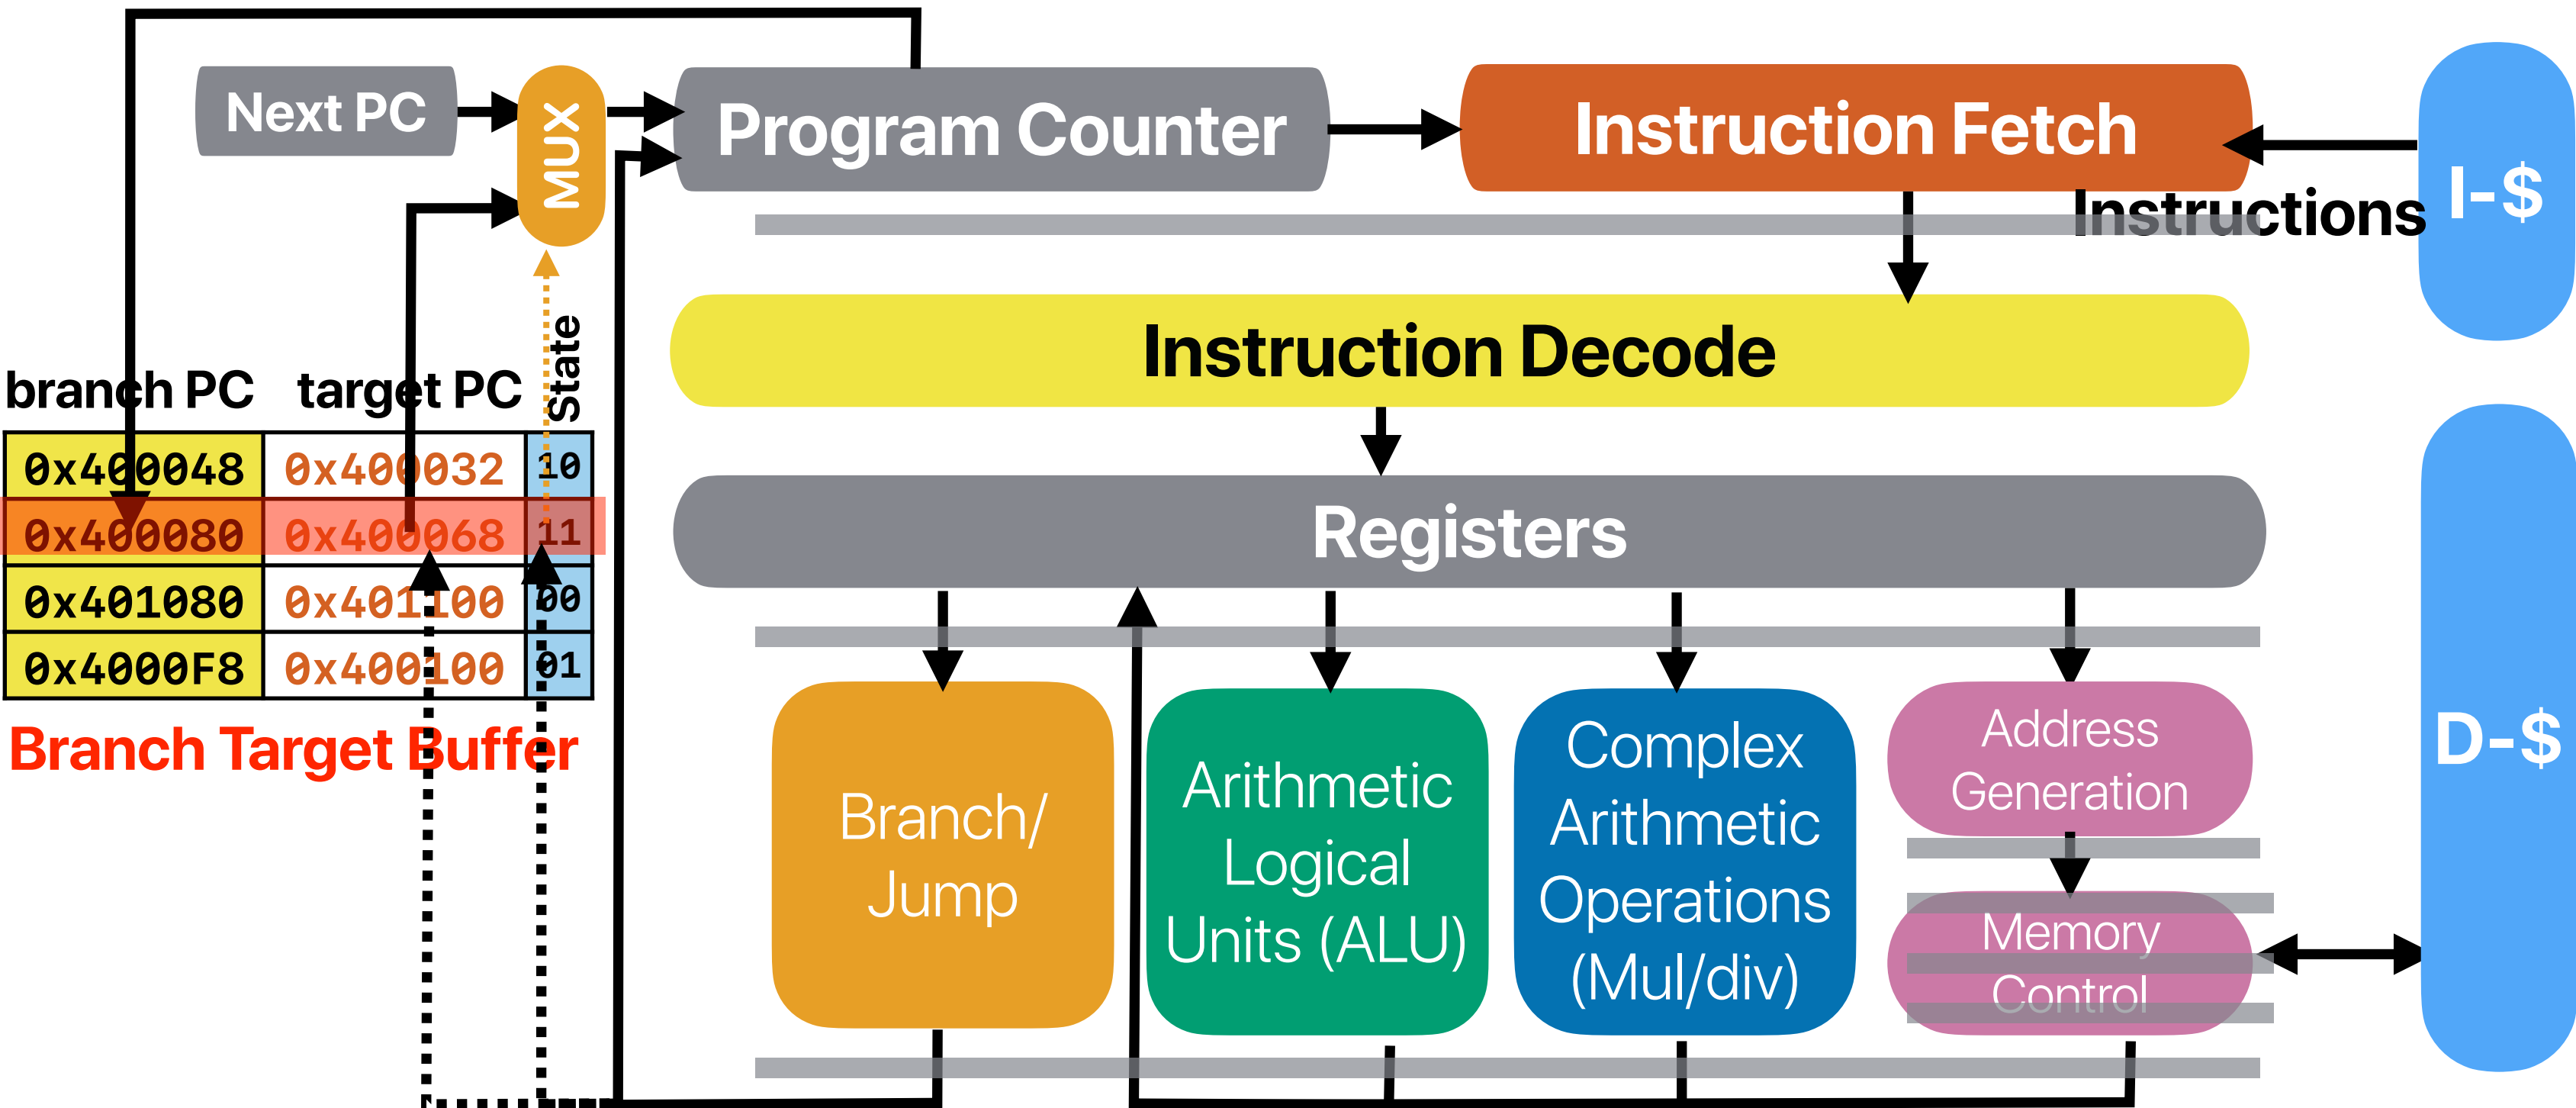
call\_when\_false(&data[i]);

branch taken (changing address to L11) if the evaluation is not true

In if-else statement, true can mean "not taken"

```
.L12:
    movl    -32(%rbp), %eax
    cltq
    leaq    0(,%rax,8), %rdx
    movq    -8(%rbp), %rax
    addq    %rdx, %rax
    movq    (%rax), %rax
    andl    $1, %eax
    testq   %rax, %rax
    je      .L10
    movl    -32(%rbp), %eax
    cltq
    leaq    0(,%rax,8), %rdx
    movq    -8(%rbp), %rax
    addq    %rdx, %rax
    movq    %rax, %rdi
    call    call_when_true
    jmp     .L11
.L10:
    movl    -32(%rbp), %eax
    cltq
    leaq    0(,%rax,8), %rdx
    movq    -8(%rbp), %rax
    addq    %rdx, %rax
    movq    %rax, %rdi
    call    call_when_false
.L11:
    addl    $1, -32(%rbp)
.L9:
    movl    -32(%rbp), %eax
    movl    -28(%rbp), %eax
    jle     .L12
    movl    $0, %eax
    leave
```

# Detail of a basic dynamic branch predictor



# 2-bit local predictor

- What's the overall branch prediction (include both branches) accuracy for this nested for loop?

```
i = 0;
do {
    if( i % 2 != 0) // Branch X, taken if i % 2 == 0
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100) // Branch Y
```

Can we do a better job?

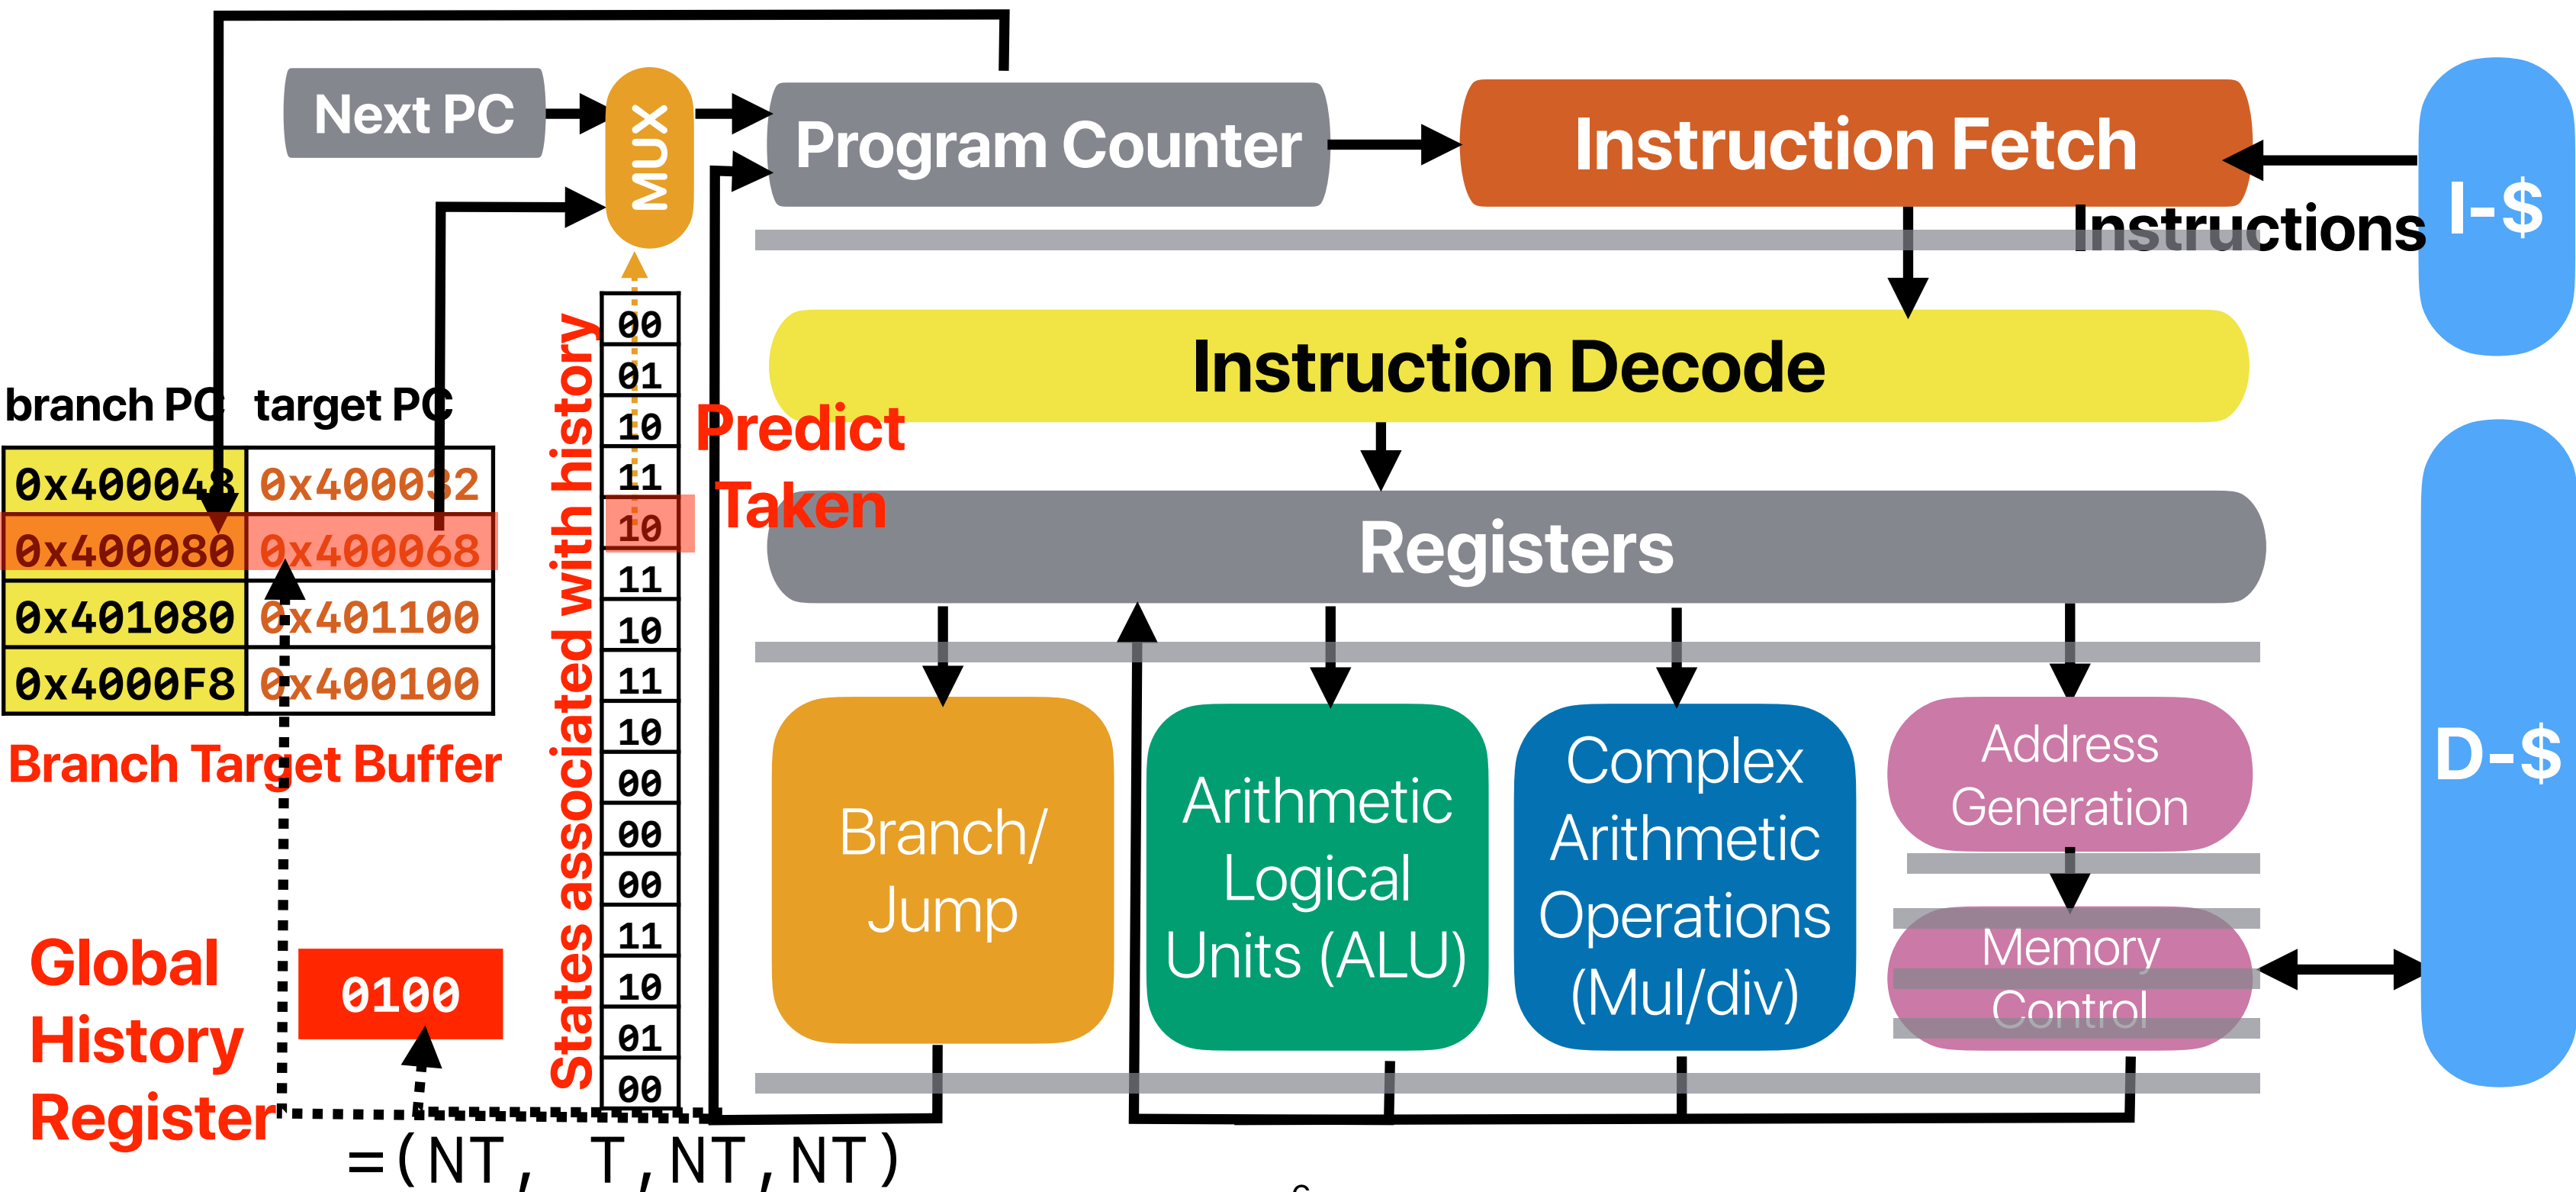
(assume all states started with 00)

- A. ~25%
- B. ~33%
- C. ~50%
- D. ~67%
- E. ~75%**

For branch Y, almost 100%,  
For branch X, only 50%

i	branch?	state	prediction	actual
0	X	00	NT	T
1	Y	00	NT	T
1	X	01	NT	NT
2	Y	01	NT	T
2	X	00	NT	T
3	Y	10	T	T
3	X	01	NT	NT
4	Y	11	T	T
4	X	00	NT	T
5	Y	11	T	T
5	X	01	NT	NT
6	Y	11	T	T
6	X	00	NT	T
7	Y	11	T	T

# Detail of a basic dynamic branch predictor



# Performance of GH predictor

```
i = 0;
do {
    if( i % 2 != 0) // Branch X, taken if i % 2 == 0
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100) // Branch Y
```

i	branch?	GHR	state	prediction	actual
0	X	000	00	NT	T
1	Y	001	00	NT	T
1	X	011	00	NT	NT
2	Y	110	00	NT	T
2	X	101	00	NT	T
3	Y	011	00	NT	T
3	X	111	00	NT	NT
4	Y	110	01	NT	T
4	X	101	01	NT	T
5	Y	011	01	NT	T
5	X	111	00	NT	NT
6	Y	110	10	T	T
6	X	101	10	T	T
7	Y	011	10	T	T
7	X	111	00	NT	NT
8	Y	110	11	T	T
8	X	101	11	T	T
9	Y	011	11	T	T
9	X	111	00	NT	NT
10	Y	110	11	T	T
10	X	101	11	T	T
11	Y	011	11	T	T

Near perfect after this



## **Demo revisited: evaluating the cost of mis-predicted branches**

- Compare the number of mis-predictions
- Calculate the difference of cycles
- We can get the “average CPI” of a mis-prediction!

**34 cycles on Intel Alder Lake**

**24 cycles on AMD Zen 3**

**Could be more expensive than cache misses**



# Better predictor?

- Consider two predictors — (L) 2-bit local predictor with unlimited BTB entries and (G) 4-bit global history with 2-bit predictors. How many of the following code snippet would allow (G) to outperform (L)?

—

```
i = 0;
do {
    if( i % 10 != 0)
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100);
```

=

```
i = 0;
do {
    a[i] += i;
} while ( ++i < 100);
```

≡

```
i = 0;
do {
    j = 0;
    do {
        sum += A[i*2+j];
    }
    while( ++j < 2);
} while ( ++i < 100);
```

≥

```
i = 0;
do {
    if( rand() %2 == 0)
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100)
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4





# Better predictor?

- Consider two predictors — (L) 2-bit local predictor with unlimited BTB entries and (G) 4-bit global history with 2-bit predictors. How many of the following code snippet would allow (G) to outperform (L)?

about the same

`i = 0;  
do {  
 if( i % 10 != 0)  
 a[i] *= 2;  
 a[i] += i;  
} while ( ++i < 100);`

about the same

`i = 0;  
do {  
 a[i] += i;  
} while ( ++i < 100);`

≡

`i = 0;  
do {  
 j = 0;  
 do {  
 sum += A[i*2+j];  
 }  
 while( ++j < 2);  
} while ( ++i < 100);`

L could be better

≥

`i = 0;  
do {  
 if( rand() %2 == 0)  
 a[i] *= 2;  
 a[i] += i;  
} while ( ++i < 100)`

A. 0

B. 1

C. 2

D. 3

E. 4

# Takeaways: branch predictions

- The cost of not to predict a branch is to stall until the data dependency is resolved — 34 cycles on modern intel processors and 24 on AMD processors
- Branch predictions allow the processor to at least make some progress and hide the stalls if we guessed correctly!
- Dynamic branch prediction — predict based on prior history
  - Local predictor — make predictions based on the state of each branch instruction
  - Global predictor — make predictions based on the state from all branches
  - Both are not perfect

# Takeaways: branch predictions

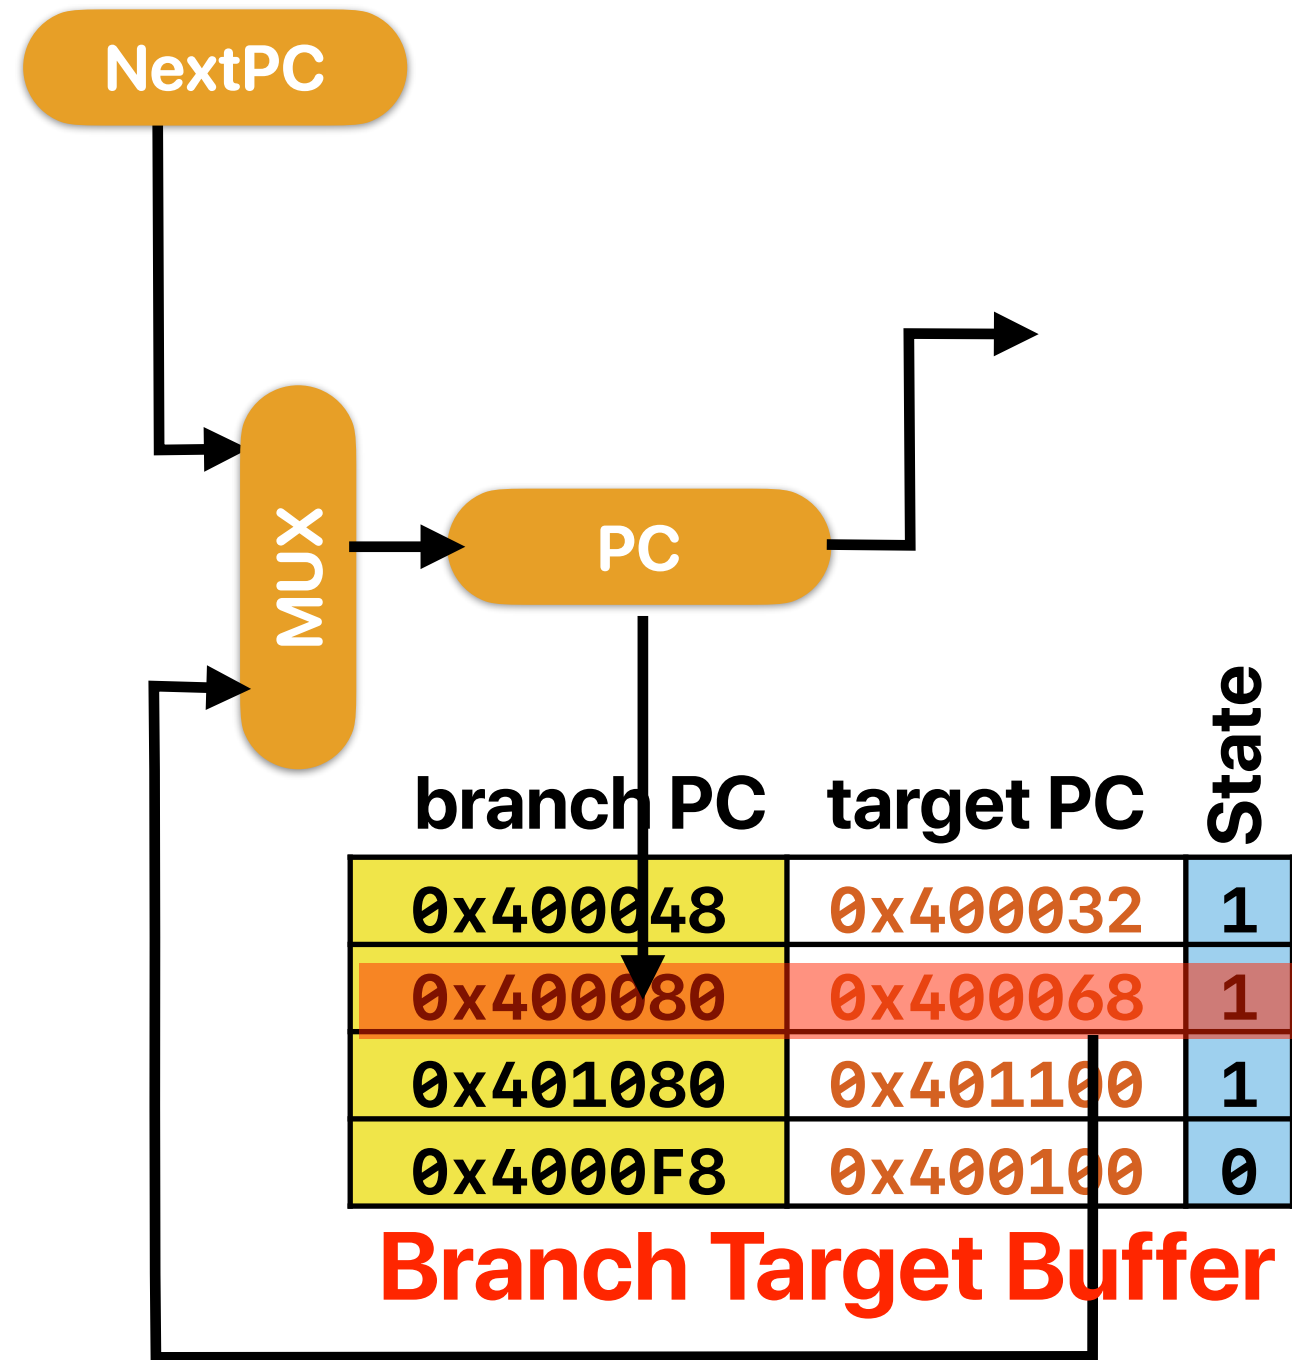
- The cost of not to predict a branch is to stall until the data dependency is resolved — 34 cycles on modern intel processors and 23 on AMD processors
- Branch predictions allow the processor to at least make some progress and hide the stalls if we guessed correctly!
- Dynamic branch prediction — predict based on prior history
  - Local predictor — make predictions based on the state of each branch instruction
  - Global predictor — make predictions based on the state from all branches
  - Both are not perfect

# Outline

- Hybrid predictors (cont.)
- Data hazards
- Hardware optimizations for data hazards

# Hybrid predictors

# Tournament Predictor



**Global History Register**

0100

States associated with history

00
01
10
11
10
11
10
11
10
11
10
00
00
00
00
00
11
10
10
01
00

**Local History Predictor**

branch PC local history

0x400048	1000
0x400080	0110
0x401080	1010
0x4000F8	0110

**Predict Taken**

States associated with history

00
01
10
11
10
11
10
11
10
11
10
00
00
00
00
00
11
10
10
01
00

# Tournament Predictor

- The state predicts “which predictor is better”
  - Local history
  - Global history
- The predicted predictor makes the prediction
- Tournament predictor is a “hybrid predictor” as it takes both local & global information into account

# TAGE

André Seznec. The L-TAGE branch predictor. Journal of Instruction Level Parallelism (<http://www.jilp.org/vol9>), May 2007.



# Better predictor?

- Consider two predictors — (L) 2-bit local predictor with unlimited BTB entries and (G) 4-bit global history with 2-bit predictors. How many of the following code snippet would allow (G) to outperform (L)?

about the same

```
i = 0;
do {
    if( i % 10 != 0)
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100);
```

about the same

```
i = 0;
do {
    a[i] += i;
} while ( ++i < 100);
```

≡

```
i = 0;
do {
    j = 0;
    do {
        sum += A[i*2+j];
    }
    while( ++j < 2);
} while ( ++i < 100);
```

L could be better

```
i = 0;
do {
    if( rand() %2 == 0)
        a[i] *= 2;
    a[i] += i;
} while ( ++i < 100)
```

A. 0

B. 1

C. 2

D. 3

E. 4

different branch needs different length of history

global predictor can work if the history is long enough!

NextPC

# TAGE

"Very" Long  
Global History  
Register

$L(N)$  — the last  $m$ -bits of  
history used for table  $N$

.....000001110100

PC

MUX

branch PC    target PC

0x400048	0x400032
0x400080	0x400068
0x401080	0x401100
0x4000F8	0x400100

Branch Target Buffer

Base predictor

01  
00  
10  
11  
10  
10  
11  
00  
00  
11  
10  
11  
01  
00  
00  
11

pred tag u

=?

pred tag u

=?

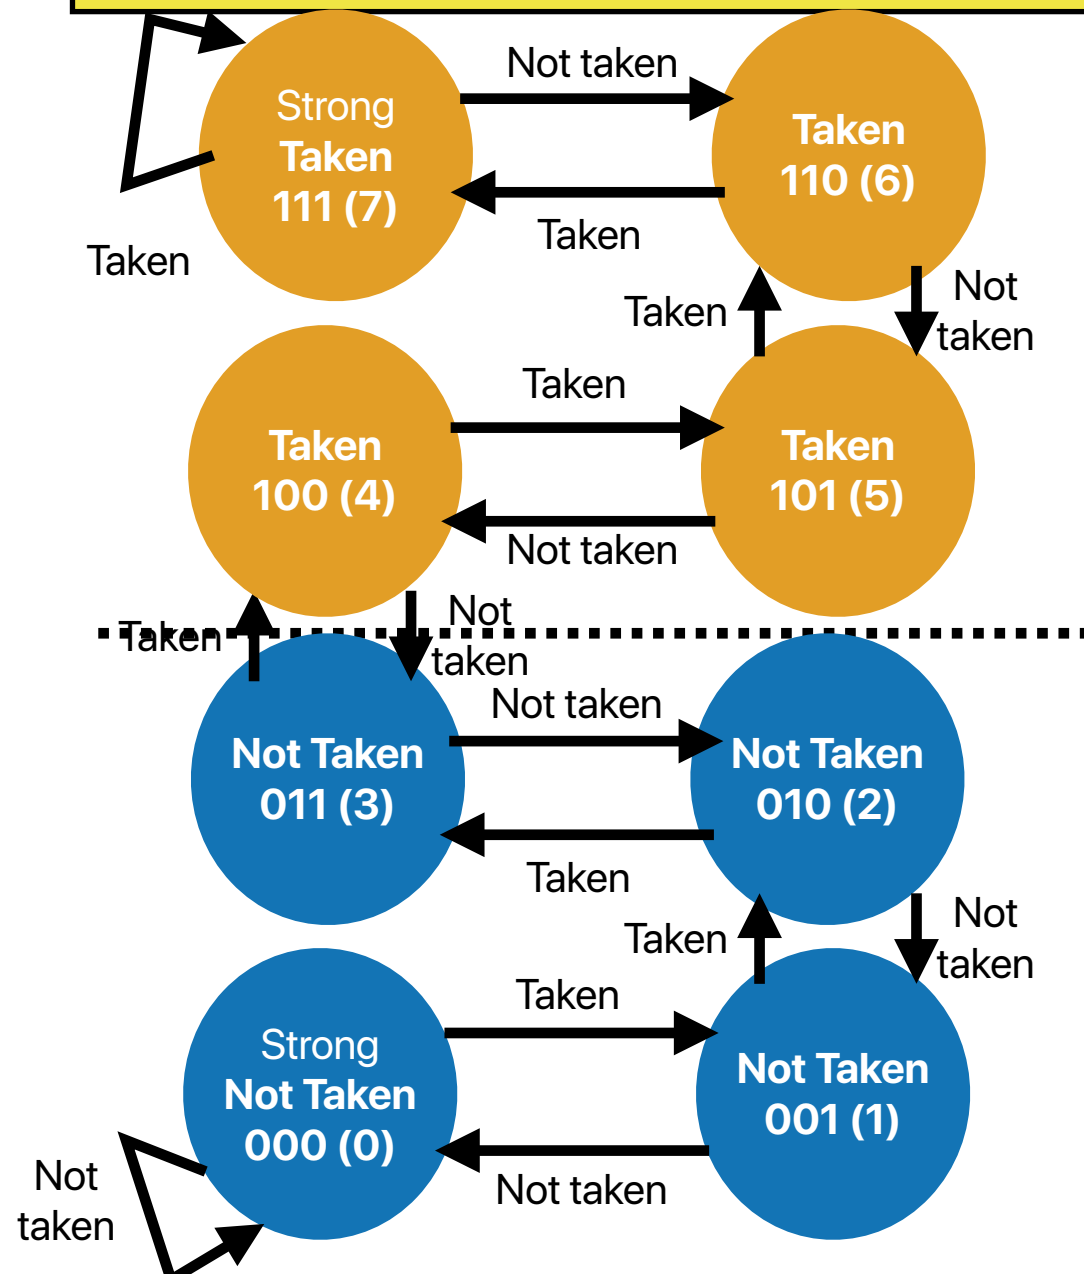
pred tag u

=?

prediction (using the  
longest match)

# What's inside each table?

<b>pred</b> <b>(3-bit counter)</b>	<b>tag</b> <b>(partial branch PC)</b>	<b>u (usefulness)</b>
---------------------------------------	--	-----------------------



*if  $prediction(alt\_predictor) \neq prediction(pred)$  :*

*if  $prediction(pred) = actual\ result$  :  $u = u + 1$*

*if  $prediction(pred) \neq actual\ result$  :  $u = u - 1$*

NextPC

# TAGE

"Very" Long  
Global History  
Register

$L(N)$  — the last  $m$ -bits of  
history used for table  $N$

.....000001110100

branch PC	target PC
0x400048	0x400032
0x400080	0x400068
0x401080	0x401100
0x4000F8	0x400100

Branch Target Buffer

Base predictor  
is used if there  
is no match

Base predictor

01  
00  
10  
11  
10  
10  
11  
00  
00  
11  
10  
11  
01  
00  
00  
11

pred  
001

tag  
0x0080

u  
1

=?

pred

tag

u

=?

pred  
110

tag  
0x0080

u  
1

=?

The longest  
match is used  
for prediction

# Perceptron

Jiménez, Daniel, and Calvin Lin. "Dynamic branch prediction with perceptrons." Proceedings HPCA Seventh International Symposium on High-Performance Computer Architecture. IEEE, 2001.

The following slides are excerpted from <https://www.jilp.org/cbp/Daniel-slides.PDF> by Daniel Jiménez

# Branch Prediction is Essentially an ML Problem

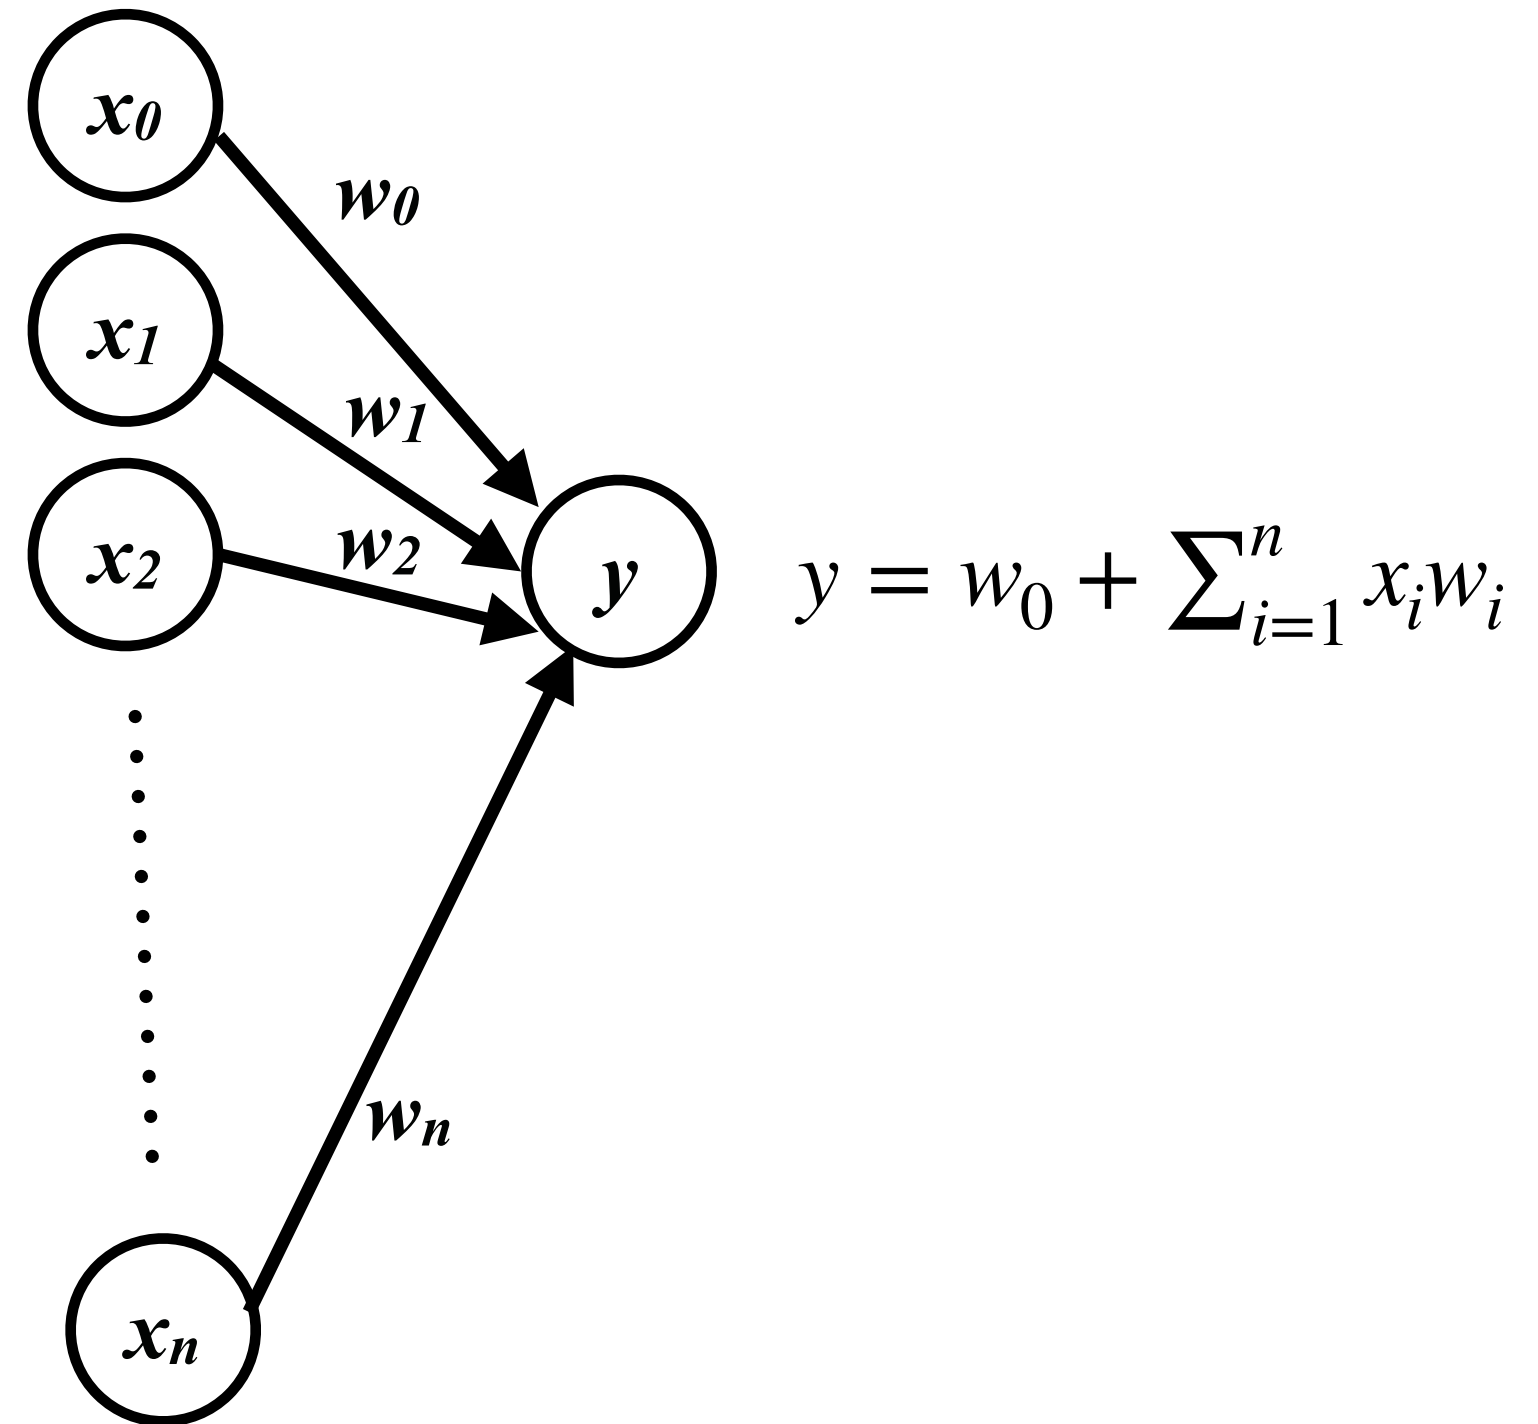
- The machine learns to predict conditional branches
- Artificial neural networks
  - Simple model of neural networks in brain cells
  - Learn to recognize and classify patterns

# Mapping Branch Prediction to NN

- The inputs to the perceptron are branch outcome histories
  - Just like in 2-level adaptive branch prediction
  - Can be global or local (per-branch) or both (alloyed)
  - Conceptually, branch outcomes are represented as
    - +1, for taken
    - -1, for not taken
- The output of the perceptron is
  - Non-negative, if the branch is predicted taken
  - Negative, if the branch is predicted not taken
- Ideally, each static branch is allocated its own perceptron

# Mapping Branch Prediction to NN (cont.)

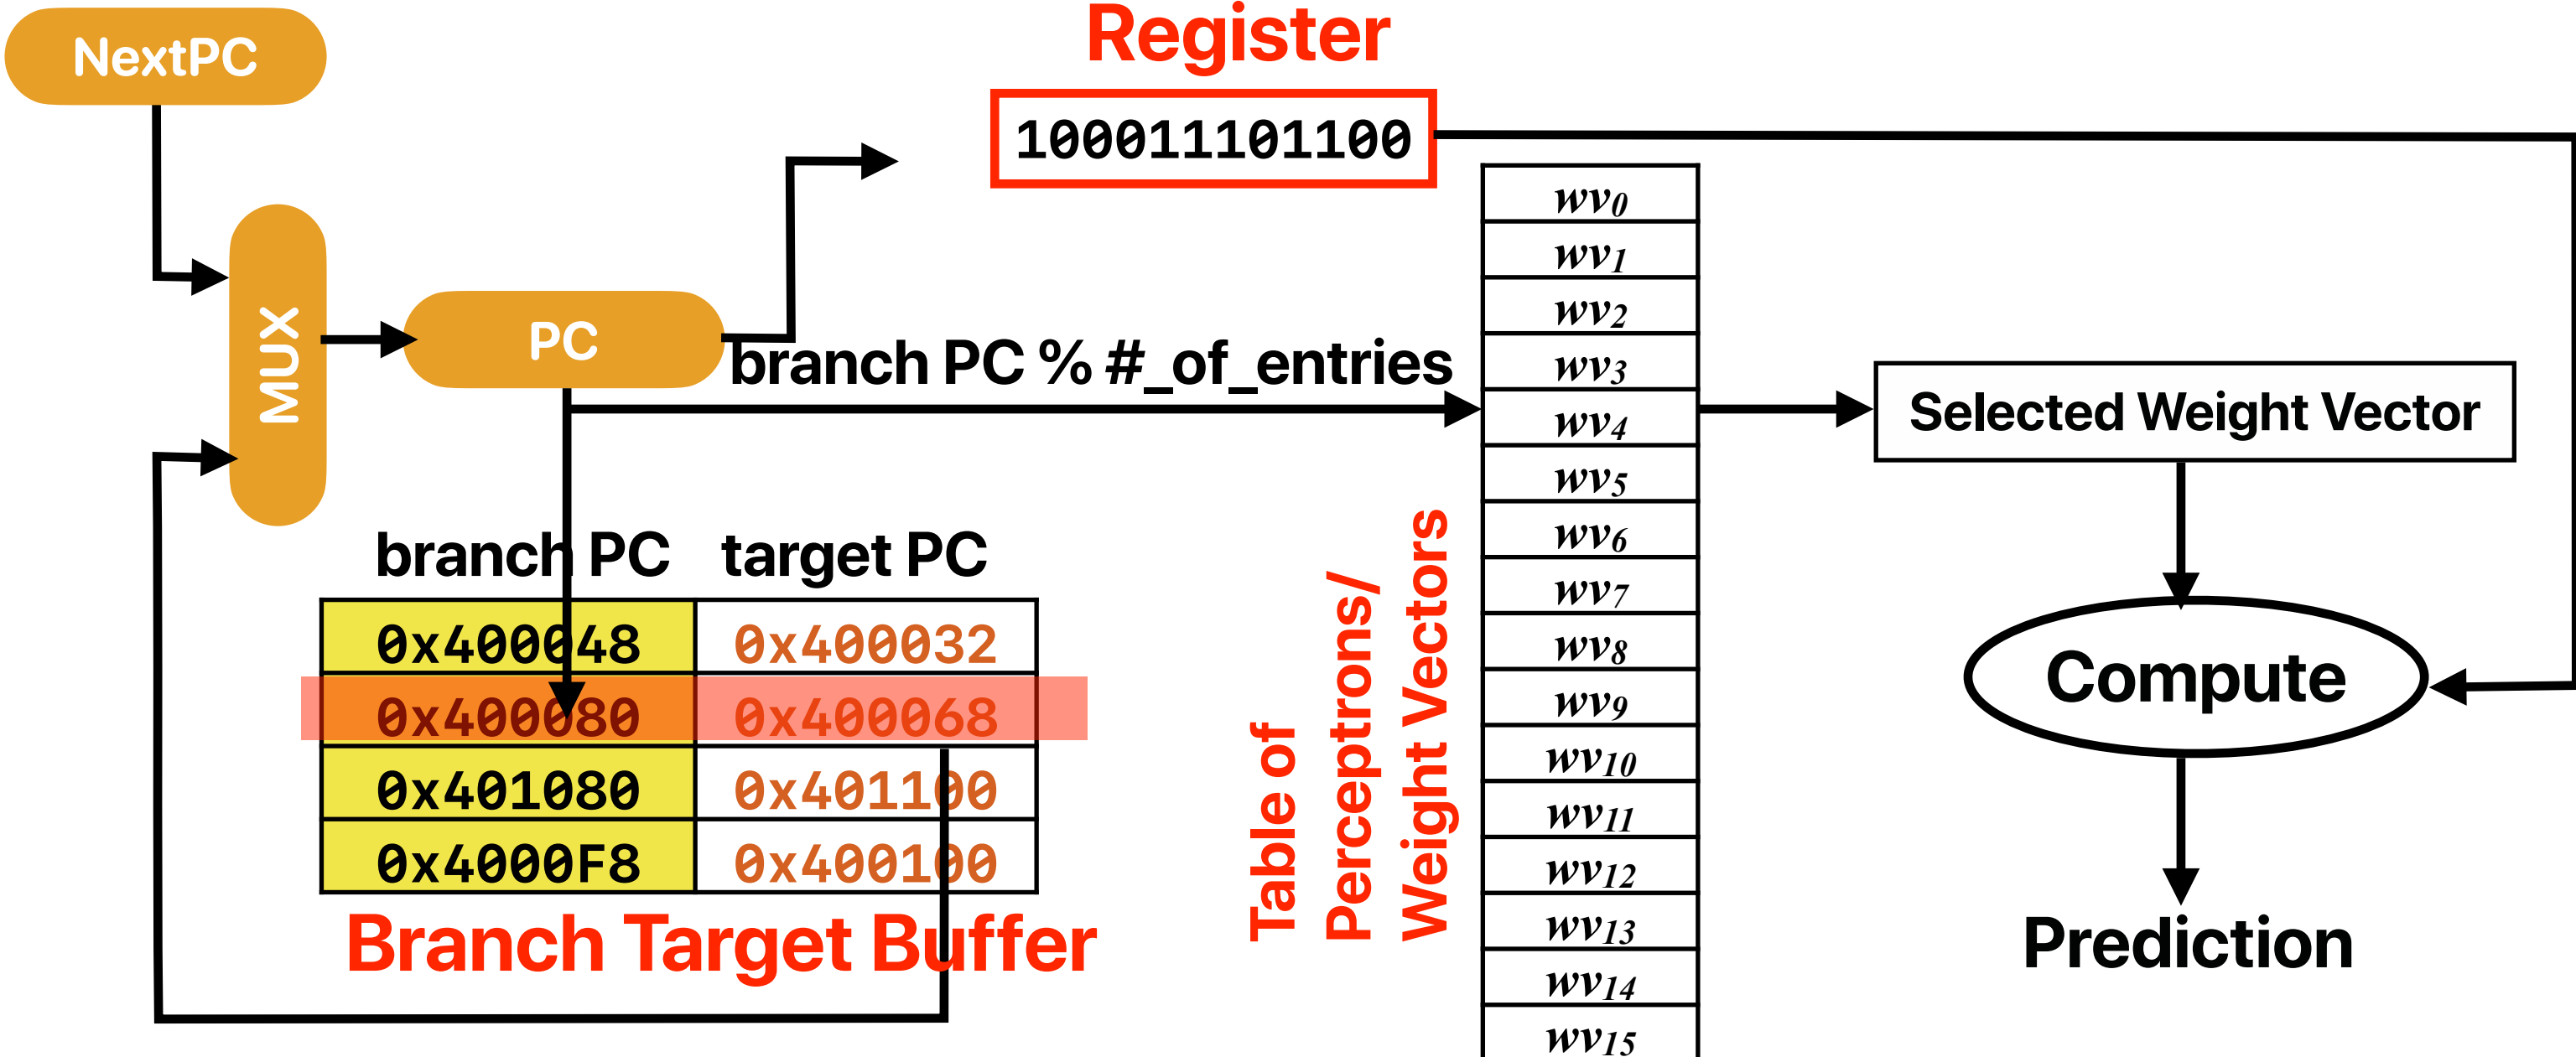
- Inputs ( $x$ 's) are from branch history and are -1 or +1
- $n + 1$  small integer weights ( $w$ 's) learned by on-line training
- Output ( $y$ ) is dot product of  $x$ 's and  $w$ 's; predict taken if  $y = 0$
- Training finds correlations between history and outcome





# Predictor Organization

Global  
History  
Register



# Training Algorithm


$x_{1..n}$  is the  $n$ -bit history register,  $x_0$  is 1.  
 $w_{0..n}$  is the weights vector.  
 $t$  is the Boolean branch outcome.  
 $\theta$  is the training threshold.

```
if  $|y| \leq \theta$  or  $((y \geq 0) \neq t)$  then
  for each  $0 \leq i \leq n$  in parallel
    if  $t = x_i$  then
       $w_i := w_i + 1$ 
    else
       $w_i := w_i - 1$ 
    end if
  end for
end if
```

**Global History** 1 0 0 0 1 1 1 0 1 1 0 0

**Branch X Taken** 

**Global History** 1 1 1 0 1 1 0 0 1 0 0 0

**Branch X Taken** 

**Global History** 1 1 0 0 1 0 0 0 1 1 1 0

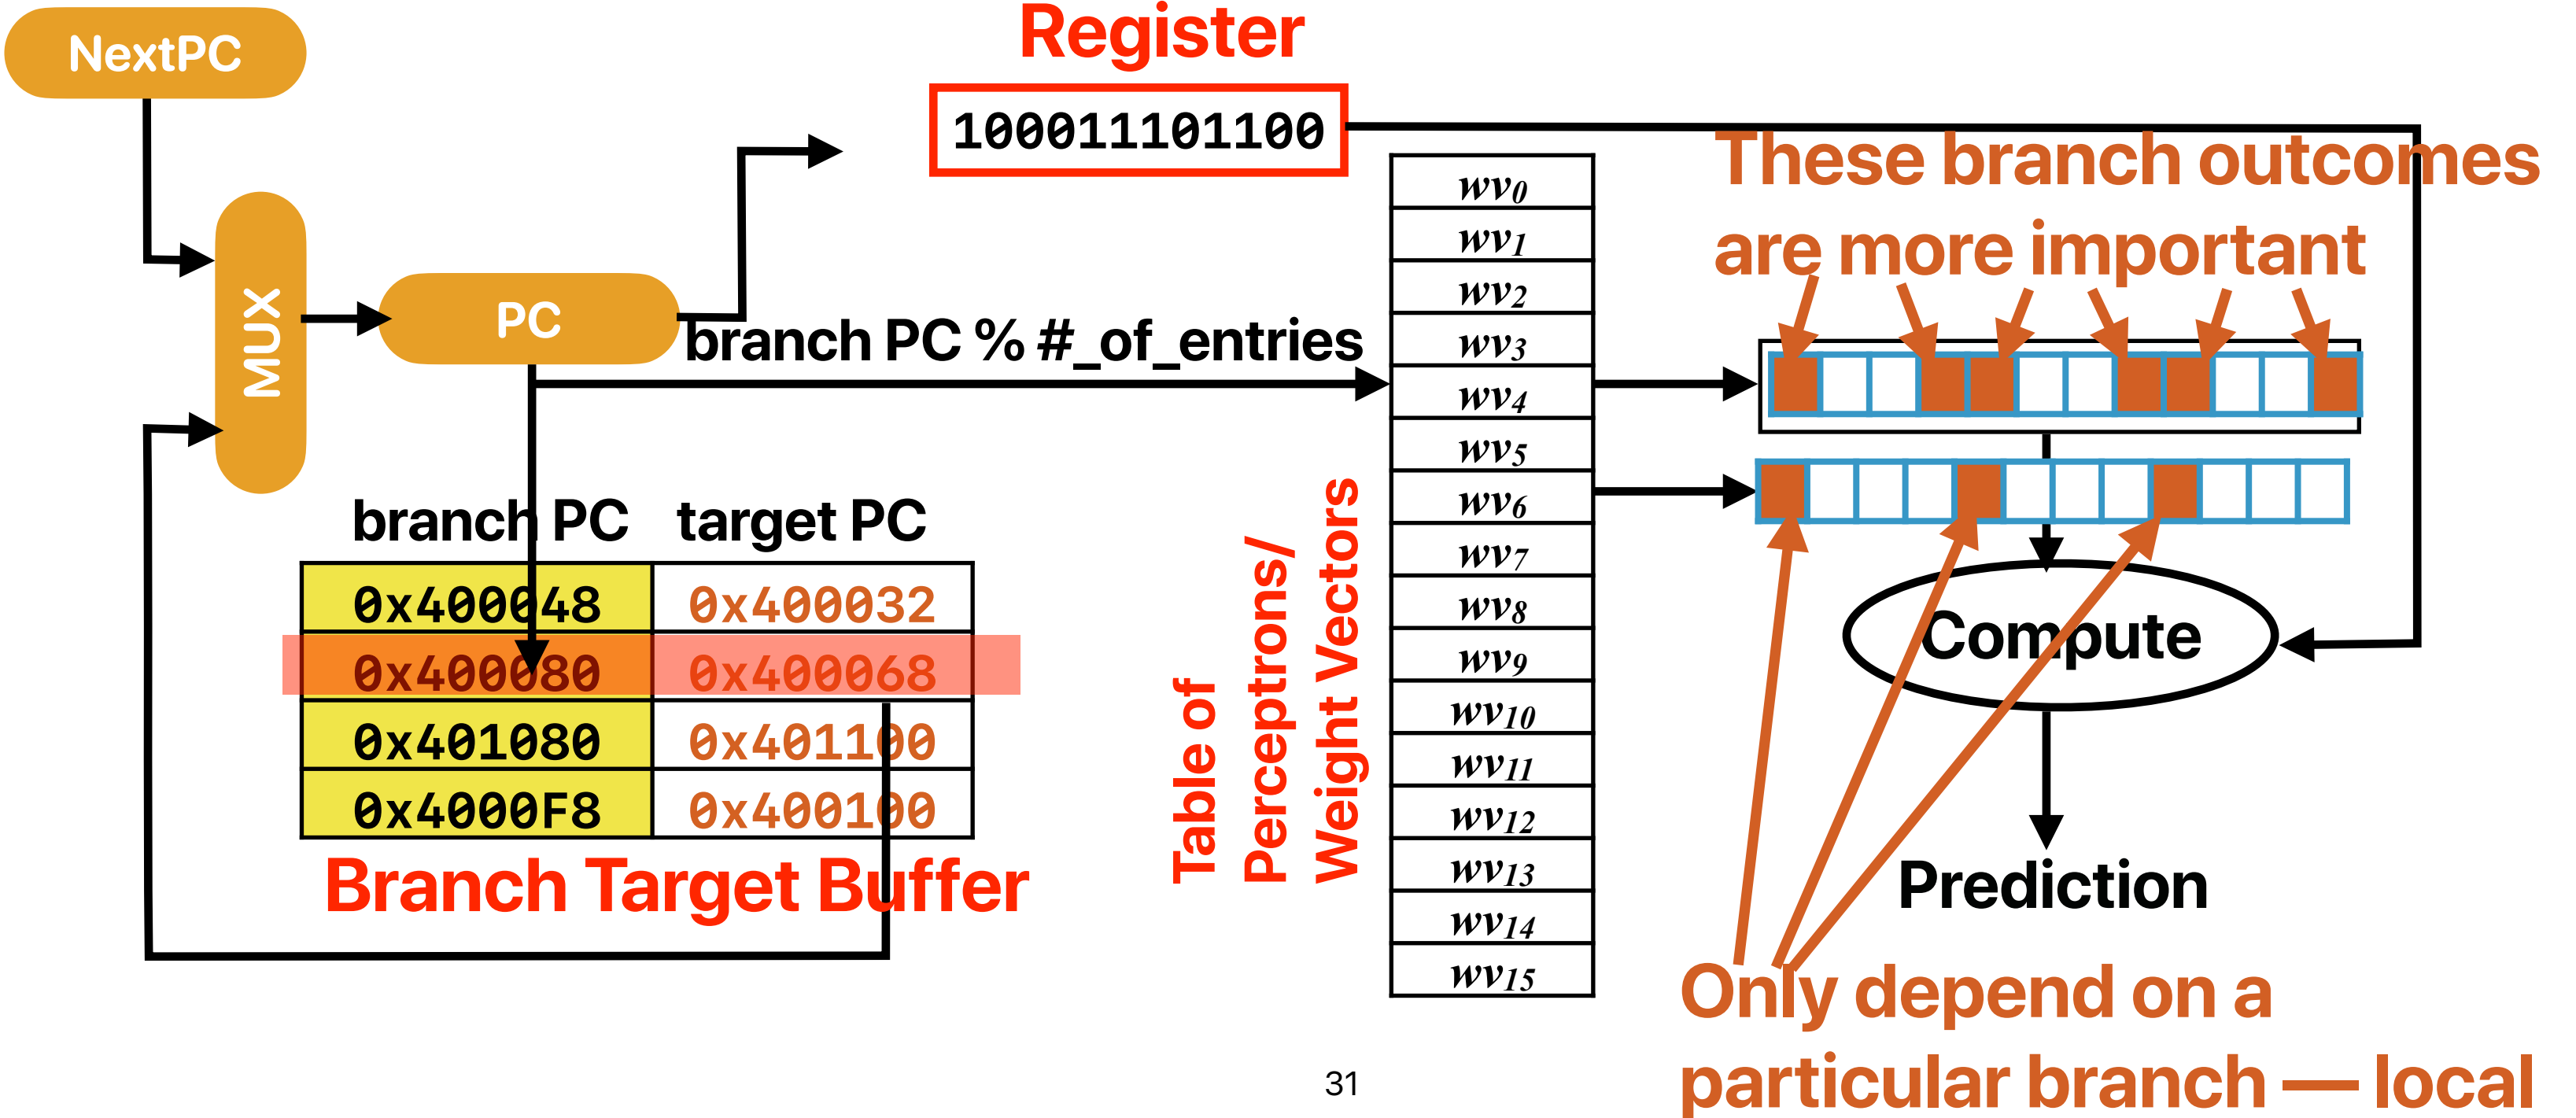
**Branch X Taken** 

**Global History** 1 0 0 0 1 1 1 0 1 1 0 0

**Branch X Taken** 

# Predictor Organization

Global  
History  
Register



# Branch predictors in processors

- The Intel Pentium MMX, Pentium II, and Pentium III have local branch predictors with a local 4-bit history and a local pattern history table with 16 entries for each conditional jump.
- Global branch prediction is used in Intel Pentium M, Core, Core 2, and Silvermont-based Atom processors.
- Tournament predictor is used in DEC Alpha, AMD Athlon processors
- The AMD Ryzen multi-core processor's Infinity Fabric and the Samsung Exynos processor include a perceptron based neural branch predictor.

# Demo revisited

**SELECT count(\*) FROM TABLE WHERE val >= A;**

**option = 1 is faster!!!, this is just a one-time cost and you don't do this after the data is loaded**

```
if(option)
    std::sort(data, data + arraySize);

for (unsigned i = 0; i < 100000; ++i) {
    int threshold = std::rand();
    for (unsigned i = 0; i < arraySize; ++i) {
        if (data[i] >= threshold)
            sum ++;
    }
}
```

	Without sorting	With sorting
The prediction accuracy of X before threshold	60%-70%	100%
The prediction accuracy of X after threshold	60%-70%	100%



# Design decisions in real practice

- AMD Zen 2 (RyZen 3000 series processors) adopts a design with first level predictor using perceptron and using TAGE for the 2<sup>nd</sup> level. What such a design decision implies about the characteristics of TAGE and Perceptron?

- ① Perceptron takes longer to train than TAGE
- ② Perceptron takes longer to predict than TAGE
- ③ Perceptron is more accurate than TAGE
- ④ Perceptron's performance improves less given more area

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4

perceptron predictors. For this reason, TAGE was a good choice for [REDACTED] L2 predictor while keeping perceptron as the L1 predictor for [REDACTED].



# Design decisions in real practice

- AMD Zen 2 (RyZen 3000 series processors) adopts a design with first level predictor using perceptron and using TAGE for the 2<sup>nd</sup> level. What such a design decision implies about the characteristics of TAGE and Perceptron?


- ① Perceptron takes longer to train than TAGE
- ② Perceptron takes longer to predict than TAGE
- ③ Perceptron is more accurate than TAGE
- ④ Perceptron's performance improves less given more area

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4


perceptron predictors. For this reason, TAGE was a good choice for [REDACTED] L2 predictor while keeping perceptron as the L1 predictor for [REDACTED].



# Area efficiency between TAGE and Perceptron



TAGE



Perceptron



# How good is prediction using perceptrons?




Figure 4: Misprediction Rates at a 4K budget. The perceptron predictor has a lower misprediction rate than *gshare* for all benchmarks except for *186.crafty* and *197.parser*.


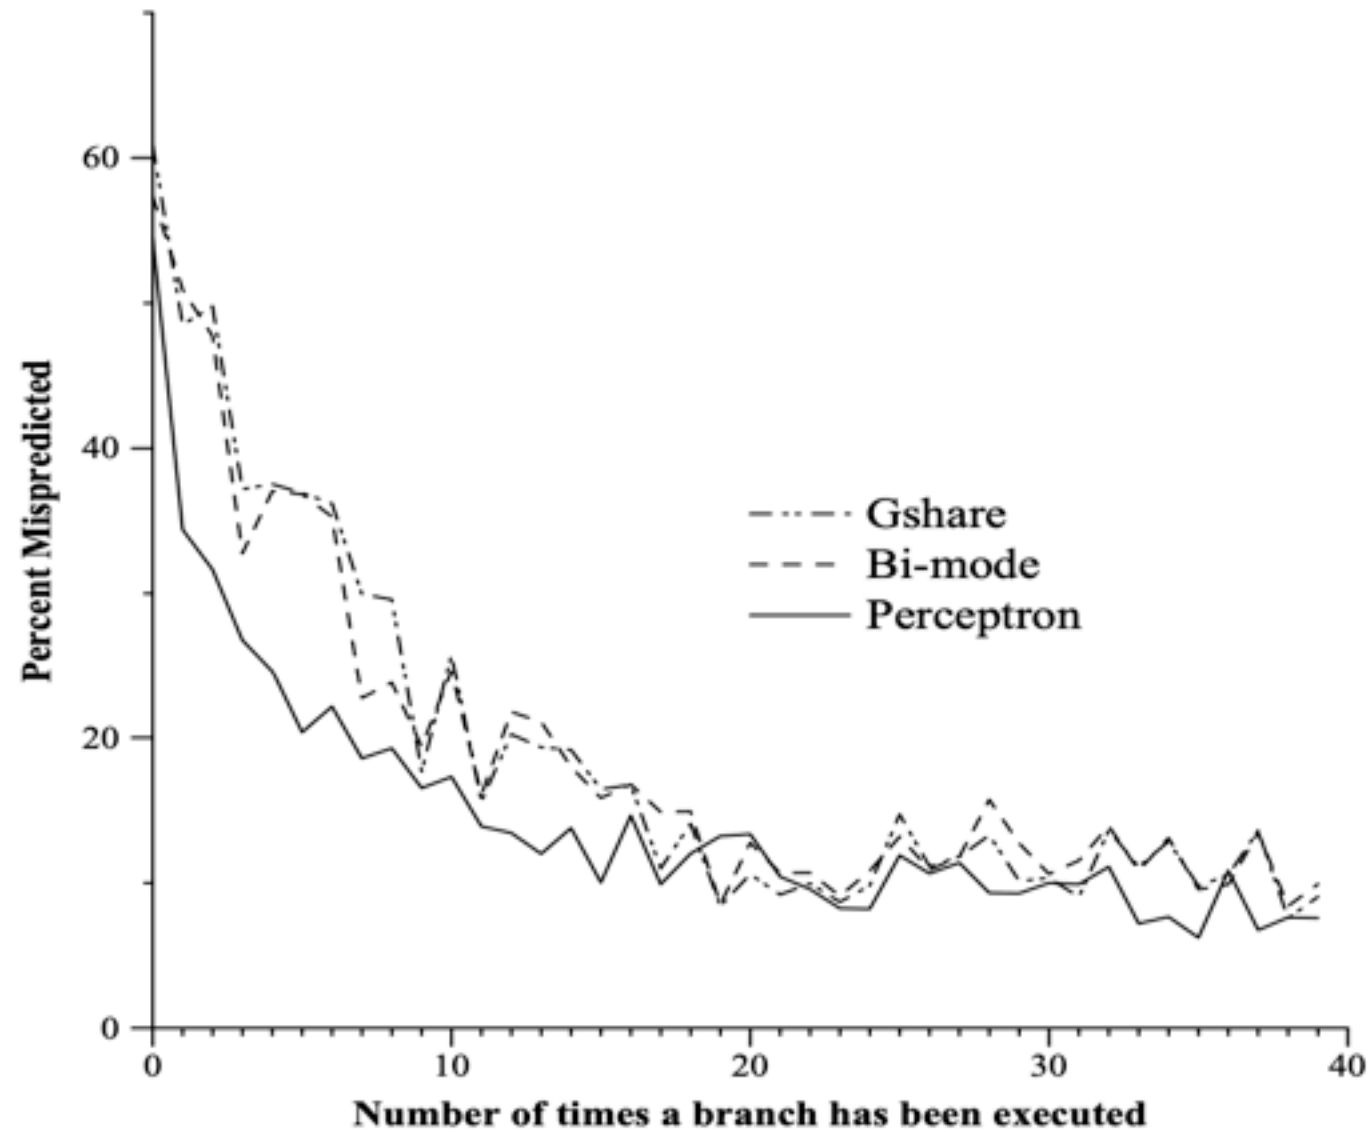


Figure 5: Misprediction Rates at a 16K budget. *Gshare* outperforms the perceptron predictor only on *186.crafty*. The hybrid predictor is consistently better than the PHT schemes.

# History/training for perceptrons



Hardware budget in kilobytes	History Length		
	<i>gshare</i>	bi-mode	perceptron
1	6	7	12
2	8	9	22
4	8	11	28
8	11	13	34
16	14	14	36
32	15	15	59
64	15	16	59
128	16	17	62
256	17	17	62
512	18	19	62

Table 1: Best History Lengths. This table shows the best amount of global history to keep for each of the branch prediction schemes.

# AMD Zen 2's design experience

## PREDICTION, FETCH, AND DECODE

The in-order front-end of the Zen 2 core includes branch prediction, instruction fetch, and decode. The branch predictor in Zen 2 features a two-level conditional branch predictor. To increase prediction accuracy, the L2 predictor has been upgraded from a perceptron predictor in Zen to a tagged geometric history length (TAGE) predictor in Zen 2.<sup>5</sup> TAGE predictors provide high accuracy per bit of storage capacity. However, they do multiplex read data from multiple tables, requiring a timing tradeoff versus perceptron predictors. For this reason, TAGE was a good choice for the longer-latency L2 predictor while keeping perceptron as the L1 predictor for best timing at low latency.

D. Suggs D. Bouvier M. Subramony and K. Lepak "Zen 2" Hot Chips vol. 31 2019.

# Design decisions in real practice

- AMD Zen 2 (RyZen 3000 series processors) adopts a design with first level predictor using perceptron and using TAGE for the 2<sup>nd</sup> level. What such a design decision implies about the characteristics of TAGE and Perceptron?

- ① Perceptron takes longer to train than TAGE  
**no — based on the paper**
- ② Perceptron takes longer to predict than TAGE  
**short — otherwise won't be in L1**
- ③ Perceptron is more accurate than TAGE  
**less accurate — otherwise won't need an L2**
- ④ Perceptron's performance improves less given more area

A. 0

**no — based on the paper**

B. 1

C. 2

D. 3

E. 4

## PREDICTION, FETCH, AND DECODE

The in-order front-end of the Zen 2 core includes branch prediction, instruction fetch, and decode. The branch predictor in Zen 2 features a two-level conditional branch predictor. To increase prediction accuracy, the L2 predictor has been upgraded from a perceptron predictor in Zen to a tagged geometric history length (TAGE) predictor in Zen 2.<sup>5</sup> TAGE predictors provide high accuracy per bit of storage capacity. However, they do multiplex read data from multiple tables, requiring a timing tradeoff versus perceptron predictors. For this reason, TAGE was a good choice for the longer-latency L2 predictor while keeping perceptron as the L1 predictor for best timing at low latency.



# Branch predictors in processors

- The Intel Pentium MMX, Pentium II, and Pentium III have local branch predictors with a local 4-bit history and a local pattern history table with 16 entries for each conditional jump.
- Global branch prediction is used in Intel Pentium M, Core, Core 2, and Silvermont-based Atom processors.
- Tournament predictor is used in DEC Alpha, AMD Athlon processors
- The AMD Ryzen multi-core processor's Infinity Fabric and the Samsung Exynos processor include a perceptron based neural branch predictor.

# Takeaways: branch predictions

- The cost of not to predict a branch is to stall until the data dependency is resolved — 34 cycles on modern intel processors and 23 on AMD processors
- Branch predictions allow the processor to at least make some progress and hide the stalls if we guessed correctly!
- Dynamic branch prediction — predict based on prior history
  - Local predictor — make predictions based on the state of each branch instruction
  - Global predictor — make predictions based on the state from all branches
  - Both are not perfect — hybrid predictors
    - Tournament
    - Perceptron
- All modern processors have pretty accurate branch predictors — if the code itself is predictable

# Recap: But A is faster!

using bit-wise operators

d. /\* one line statement using bit-wise operators \*/ (most efficient)

```
a^=b^=a^=b;
```

The order of evaluation is from right to left. This is same as in approach (c) but the three statements are compounded into one statement.

A

```
void regswap(int* a, int* b) {  
    int temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

B

```
void xorswap(int* a, int* b) {  
    *a ^= *b = *a = *b;  
}
```

# Data hazards



# Data hazards

- An instruction currently in the pipeline cannot receive the “logically” correct value for execution
- Data dependencies
  - The output of an instruction is the input of a later instruction
  - **May sometimes** result in data hazard if the later instruction that consumes the result is still in the pipeline



# How many data dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?

```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```

```
int temp = *a;
*a = *b;
*b = temp;
```

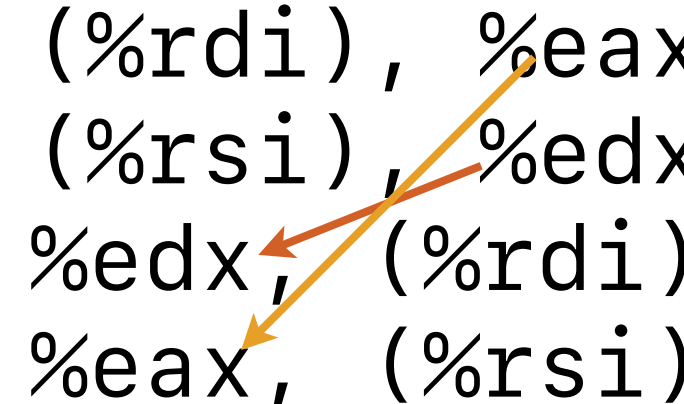
- A. 1
- B. 2
- C. 3
- D. 4
- E. 5

A screenshot of a Pollev poll interface. It shows a list of five empty input boxes, each preceded by a letter (A, B, C, D, E) in a small font. The boxes are for users to enter their answers to the question. The interface is clean and modern, with a light blue header and a white background for the poll area.

# How many dependencies do we have?

- How many pairs of data dependences are there in the following x86 instructions?

```
movl    (%rdi), %eax  
movl    (%rsi), %edx  
movl    %edx, (%rdi)  
movl    %eax, (%rsi)
```



```
int temp = *a;  
*a = *b;  
*b = temp;
```

A. 1

B. 2

C. 3

D. 4

E. 5



# How many data dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?

```
movl    (%rdi), %eax
xorl    (%rsi), %eax
movl    %eax, (%rdi)
xorl    (%rsi), %eax
movl    %eax, (%rsi)
xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5



# How many dependencies do we have?

- How many pairs of data dependencies are there in the following x86 instructions?

```
movl    (%rdi), %eax
xorl    (%rsi), %eax
movl    %eax, (%rdi)
xorl    (%rsi), %eax
movl    %eax, (%rsi)
xorl    %eax, (%rdi)
```

```
*a ^= *b;
*b ^= *a;
*a ^= *b;
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5

# Data hazards

① movl (%rdi), %eax

② movl (%rsi), %edx

③ movl %edx, (%rdi)

④ movl %eax, (%rsi)

	IF	ID	ALU/BR/AG	M1	M2	M3	M4/XORL	WB/Retire
1	(1)							
2	(2)	(1)						
3	(3)	(2)	(1)					
4	(4)	(3)	(2)	(1)				
5		(4)	(3)	(2)	(1)			
6			(4)	(3)	(2)	(1)		
7				(4)	(3)	(2)	(1)	
8					(4)	(3)	(2)	
9						(4)	(3)	
10							(4)	
11								
12								
13								
14								

%edx does not have our desired value

%eax does not have our desired value

# Solution 1: Let's try "stall" again

- Whenever the input is not ready when the consumer is decoding, just stall — the consumer stays at ID.



# Data hazards?

- How many cycles do we have to stall in the following x86 instructions to get the expected output if a memory operation (assume 100% cache hit rate) takes 5 cycles?

① `movl (%rdi), %eax`

② `movl (%rsi), %edx`

③ `movl %edx, (%rdi)`

④ `movl %eax, (%rsi)`

A. 1

B. 2

C. 3

D. 4

E. 5





# Data hazards?

- How many cycles do we have to stall in the following x86 instructions to get the expected output if a memory operation (assume 100% cache hit rate) takes 5 cycles?

```
① movl    (%rdi), %eax
```

```
② movl    (%rsi), %edx
```

```
③ movl    %edx, (%rdi)
```

```
④ movl    %eax, (%rsi)
```

- A. 1

- ## B. 2

- C. 3

- ## D. 4

- ## E. 5

	IF	ID	ALU/BR/AG	M1	M2	M3	M4/XORL	WB/Retire
1	(1)							
2	(2)	(1)						
3	(3)	(2)	(1)					
4	(4)	(3)	(2)	(1)				
5	(4)	(3)		(2)	(1)			
6	(4)	(3)			(2)	(1)		
7	(4)	(3)				(2)	(1)	
8	(4)	(3)					(2)	(1)
9	(4)	(3)						(2)
10		(4)	(3)					
11			(4)	(3)				
12				(4)	(3)			
13					(4)	(3)		
14						(4)		(3)
15								(4)

we have the value for %edx already!

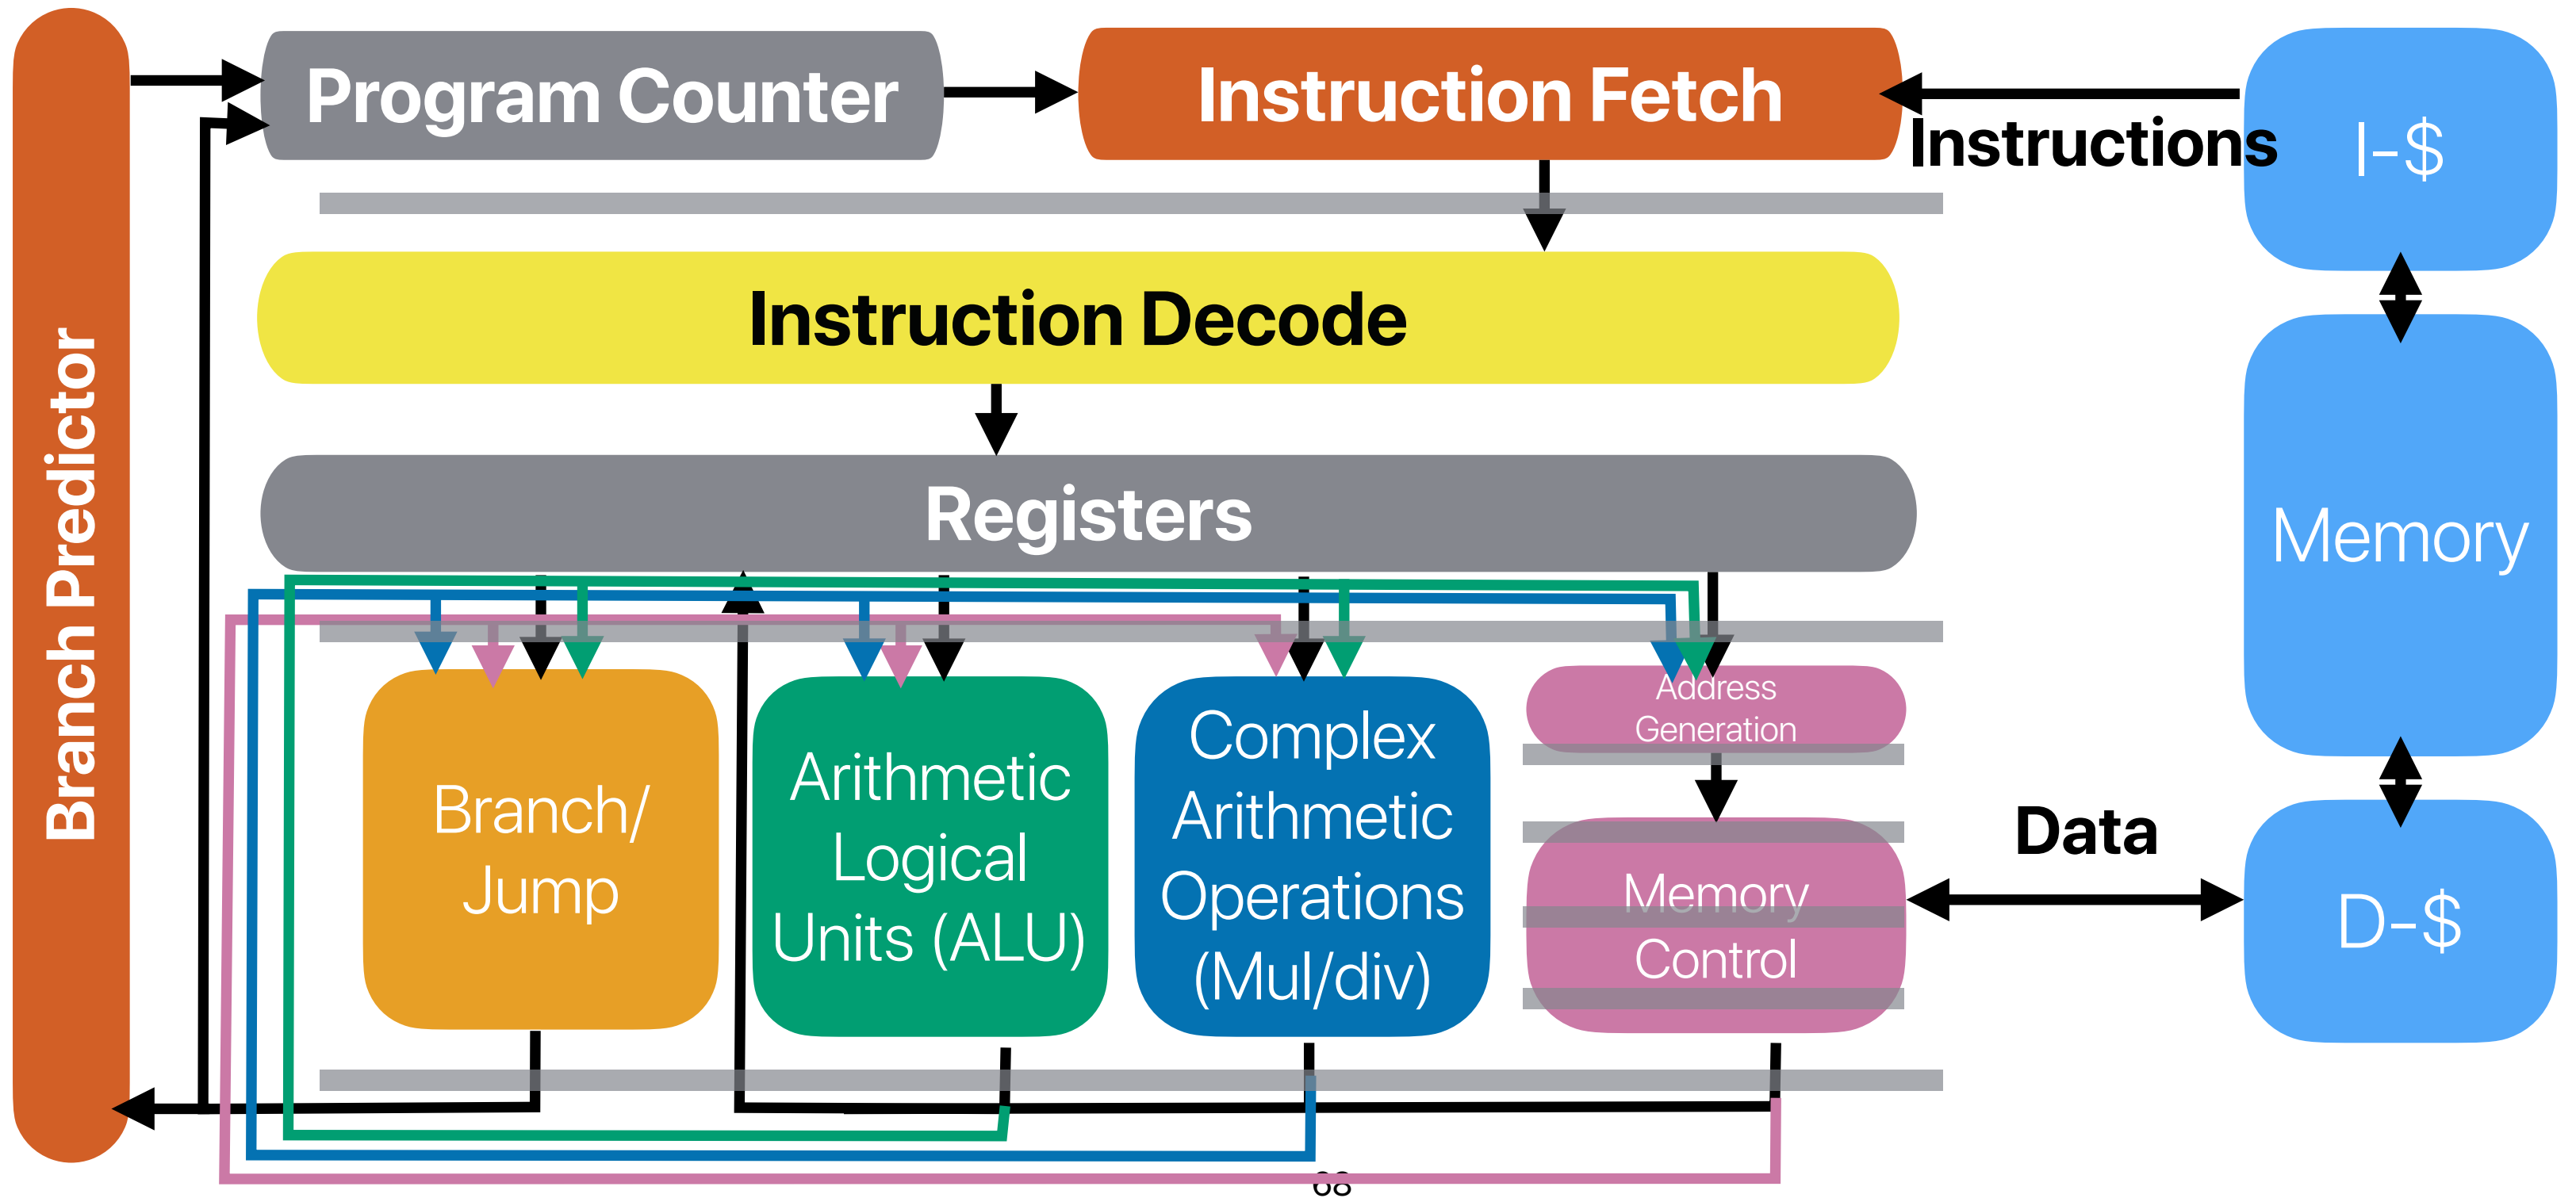
Why another cycle?

5 cycles stalls

## Solution 2: Data forwarding

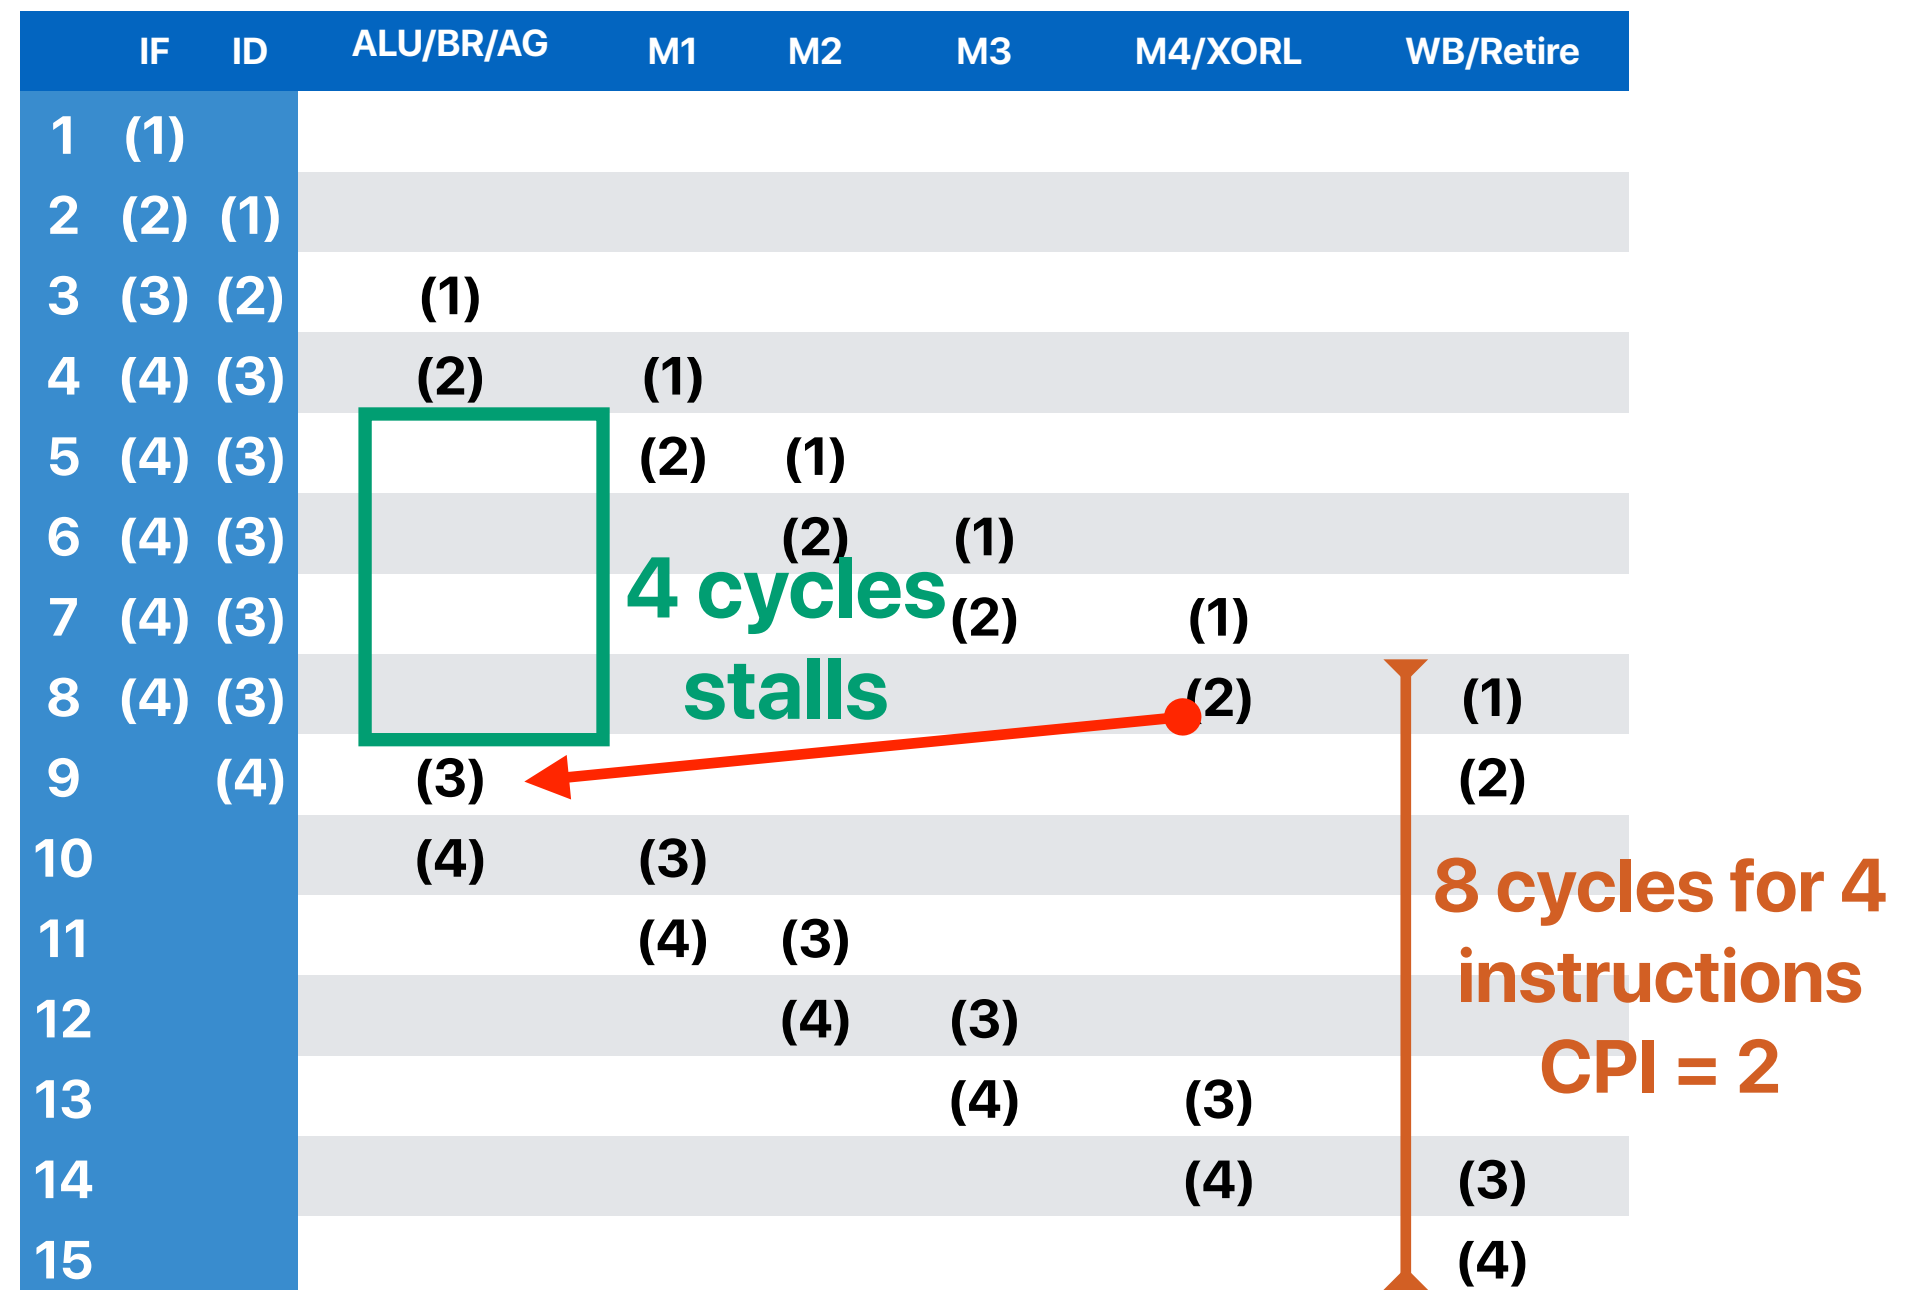
- Add logics/wires to forward the desired values to the demanding instructions

# Data "forwarding"



# The effect of data forwarding

- ① `movl (%rdi), %eax`
- ② `movl (%rsi), %edx`
- ③ `movl %edx, (%rdi)`
- ④ `movl %eax, (%rsi)`



## ★ AI Overview

Here are some tips for doing well in a technical interview:

- **Research the company:** Learn about the company's values, goals, and any problems that your skills can help with. 🔗
- **Review the job posting:** Make a list of the skills, tools, and programs required or recommended for the job. 🔗
- **Practice explaining your thought process:** When practicing coding problems explain your techniques and thought process as you work through the solution. 🔗
- **Ask clarifying questions:** Interviewers often don't provide the full picture, so ask questions to get more information. 🔗
- **Focus on the interviewer:** In an online interview, look at the webcam instead of your video image. 🔗
- **Be yourself:** Share your personality and thought processes. 🔗
- **Bring your resume:** Bring a few copies of your resume to demonstrate preparedness. 🔗
- **Master your programming language:** Focus on the assignment instead of figuring out syntax. 🔗
- **Take your time:** Don't rush yourself, and take notes if you like. 🔗
- **Invite collaboration:** The interview is interactive, and the interview team is ready to help you work through problems. 🔗

## Practice in your assignments and examines!

- Why do I pick/support this solution?
  - Prove it delivers the right result
  - Justify it's better "something" compared against other "alternatives"
- What's the solution?
  - Describe the high-level architecture/idea of your solution
- How does it work out in the problem?

# Announcements

- **Assignment 4** due next **Thursday**
- **Reading Quiz 8** due **Tuesday** before the lecture

# Computer Science & Engineering

# 203

# つづく

